

Module 2 - Assembly

Lecture 10a: Genomics

Bioinformatics Algorithms CSC4181/6802

Most slides used are from Ben Langmead's Teaching Materials (www.langmead-lab.org/teaching-materials)

Sequencing Technology

First generation

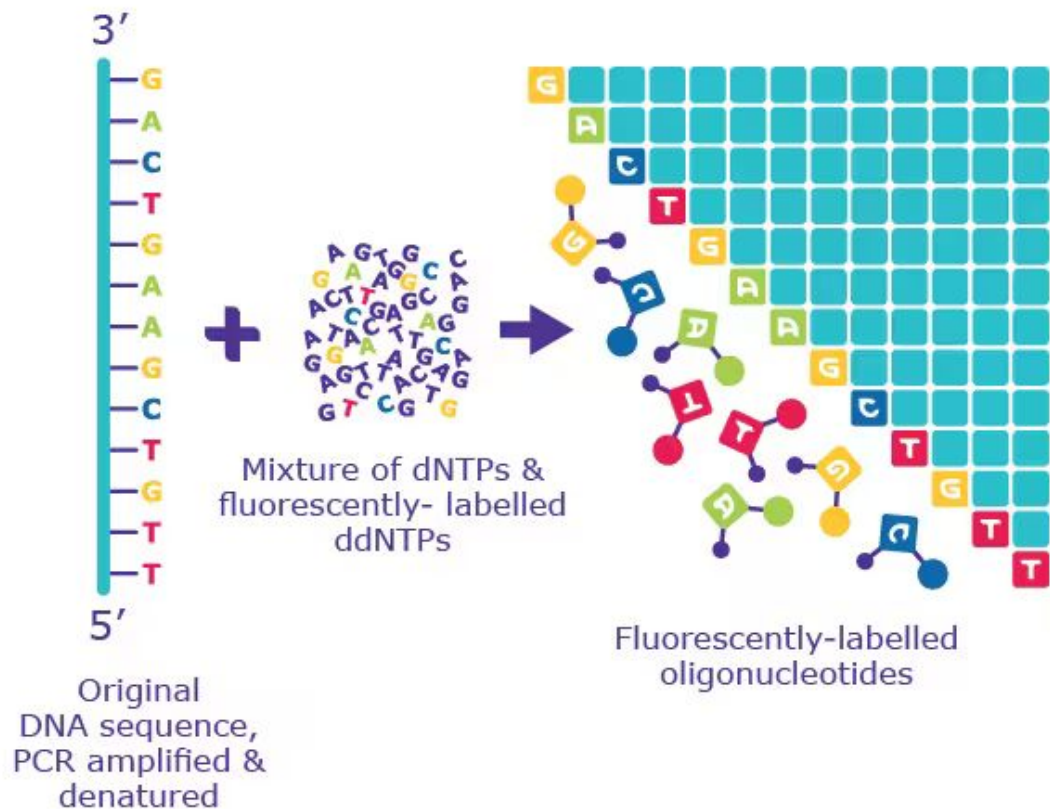


Sanger sequencing
Maxam and Gilbert
Sanger chain termination

Sanger Sequencing


1

PCR with fluorescent, chain-terminating ddNTPs




2

Size separation by capillary gel electrophoresis



3

Laser excitation & detection by sequencing machine



Sequencing Technology

First generation



Sanger sequencing
Maxam and Gilbert
Sanger chain termination

Infer nucleotide identity using dNTPs,
then visualize with electrophoresis

500–1,000 bp fragments

Sequencing Technology

First generation

Second generation
(next generation sequencing)



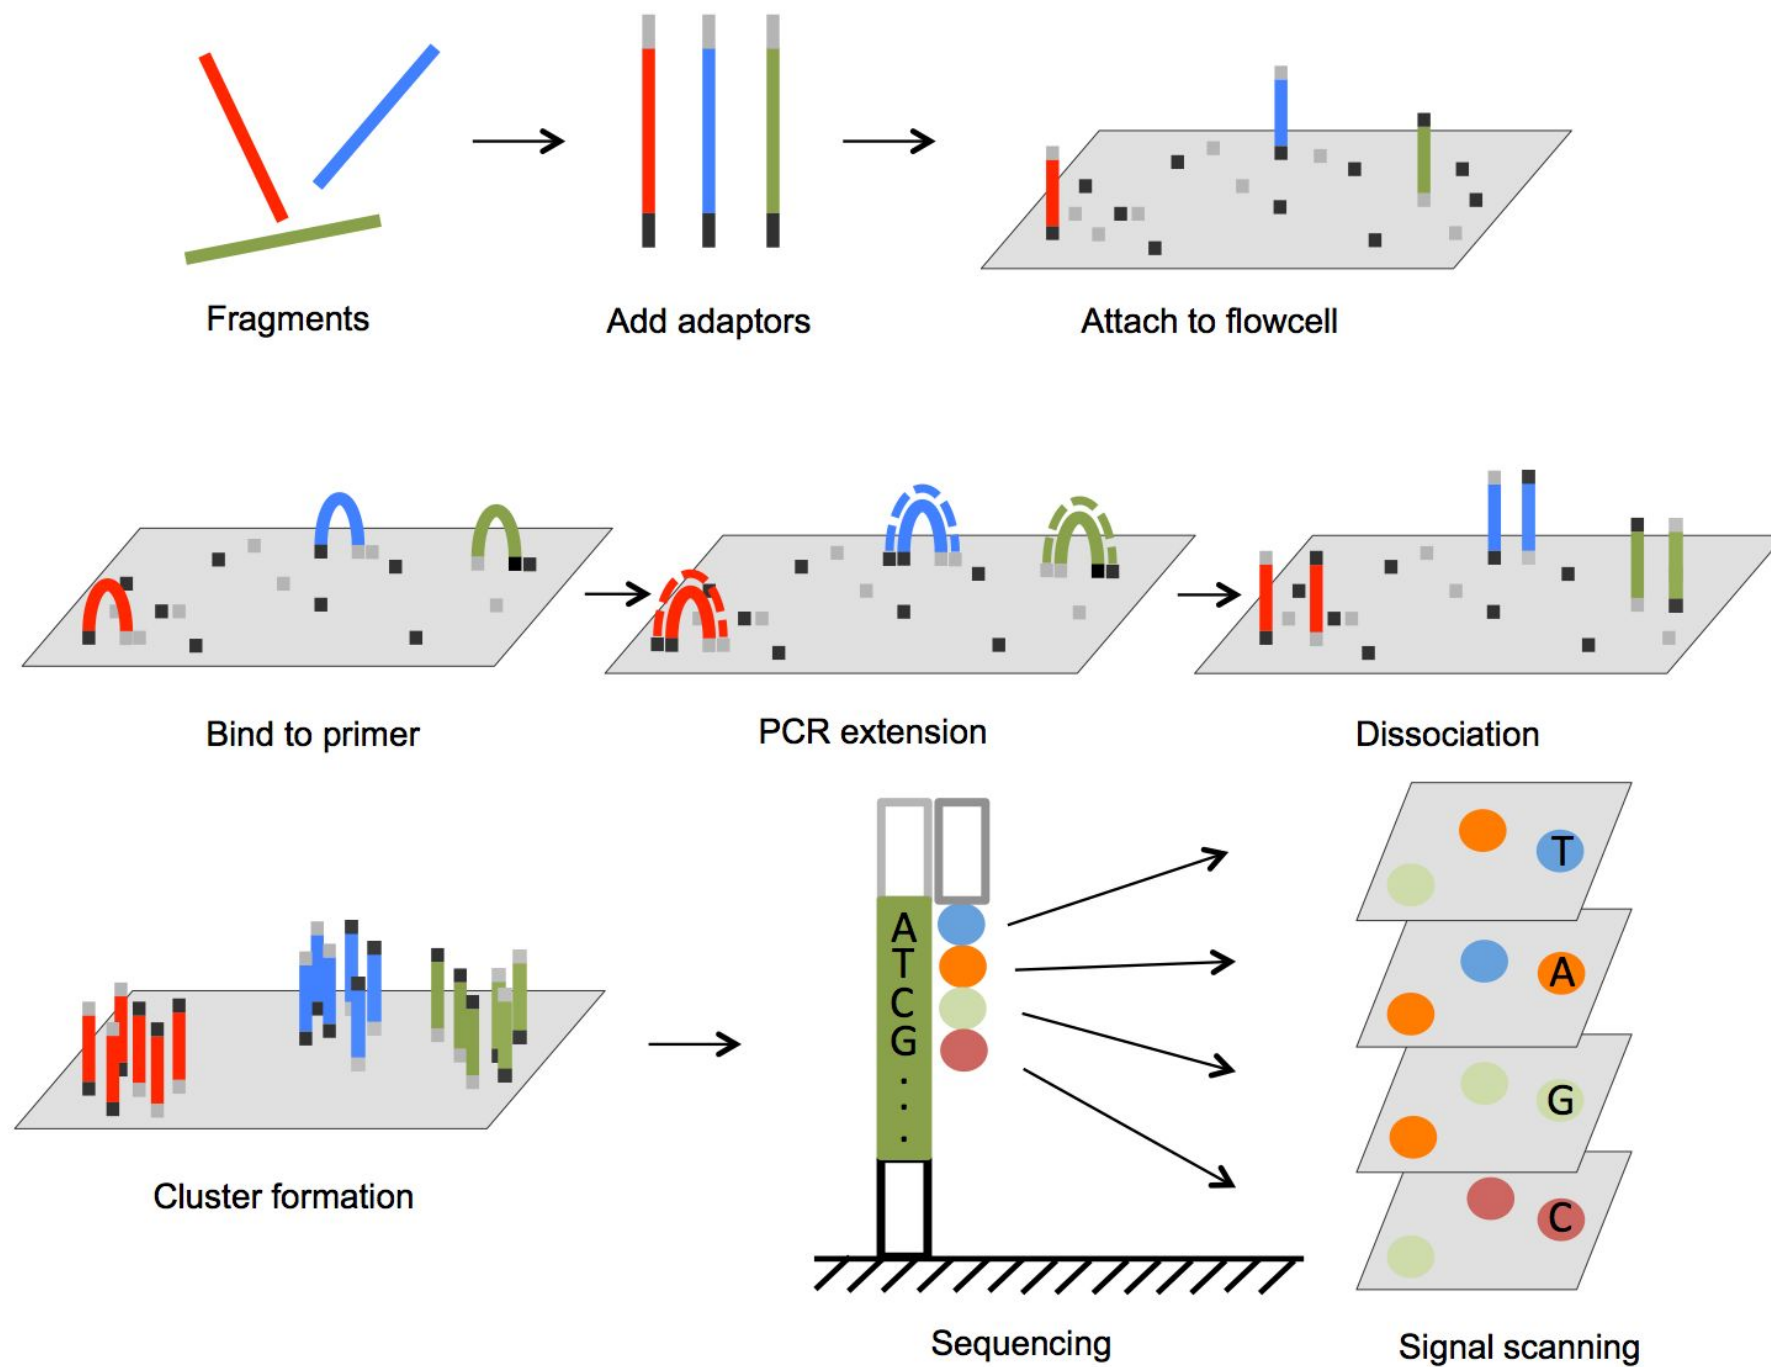
Sanger sequencing
Maxam and Gilbert
Sanger chain termination

454, Solexa,
Ion Torrent,
Illumina

Infer nucleotide identity using dNTPs,
then visualize with electrophoresis

500–1,000 bp fragments

Sequencing by Synthesis



Sequencing Technology

First generation

Second generation (next generation sequencing)



Sanger sequencing
Maxam and Gilbert
Sanger chain termination

454, Solexa,
Ion Torrent,
Illumina

Infer nucleotide identity using dNTPs,
then visualize with electrophoresis

High throughput from the
parallelization of sequencing reactions

500–1,000 bp fragments

~50–500 bp fragments

Sequencing Technology

First generation



Sanger sequencing
Maxam and Gilbert
Sanger chain termination

Infer nucleotide identity using dNTPs,
then visualize with electrophoresis

500–1,000 bp fragments

Second generation (next generation sequencing)



454, Solexa,
Ion Torrent,
Illumina

High throughput from the
parallelization of sequencing reactions


~50–500 bp fragments

Third generation



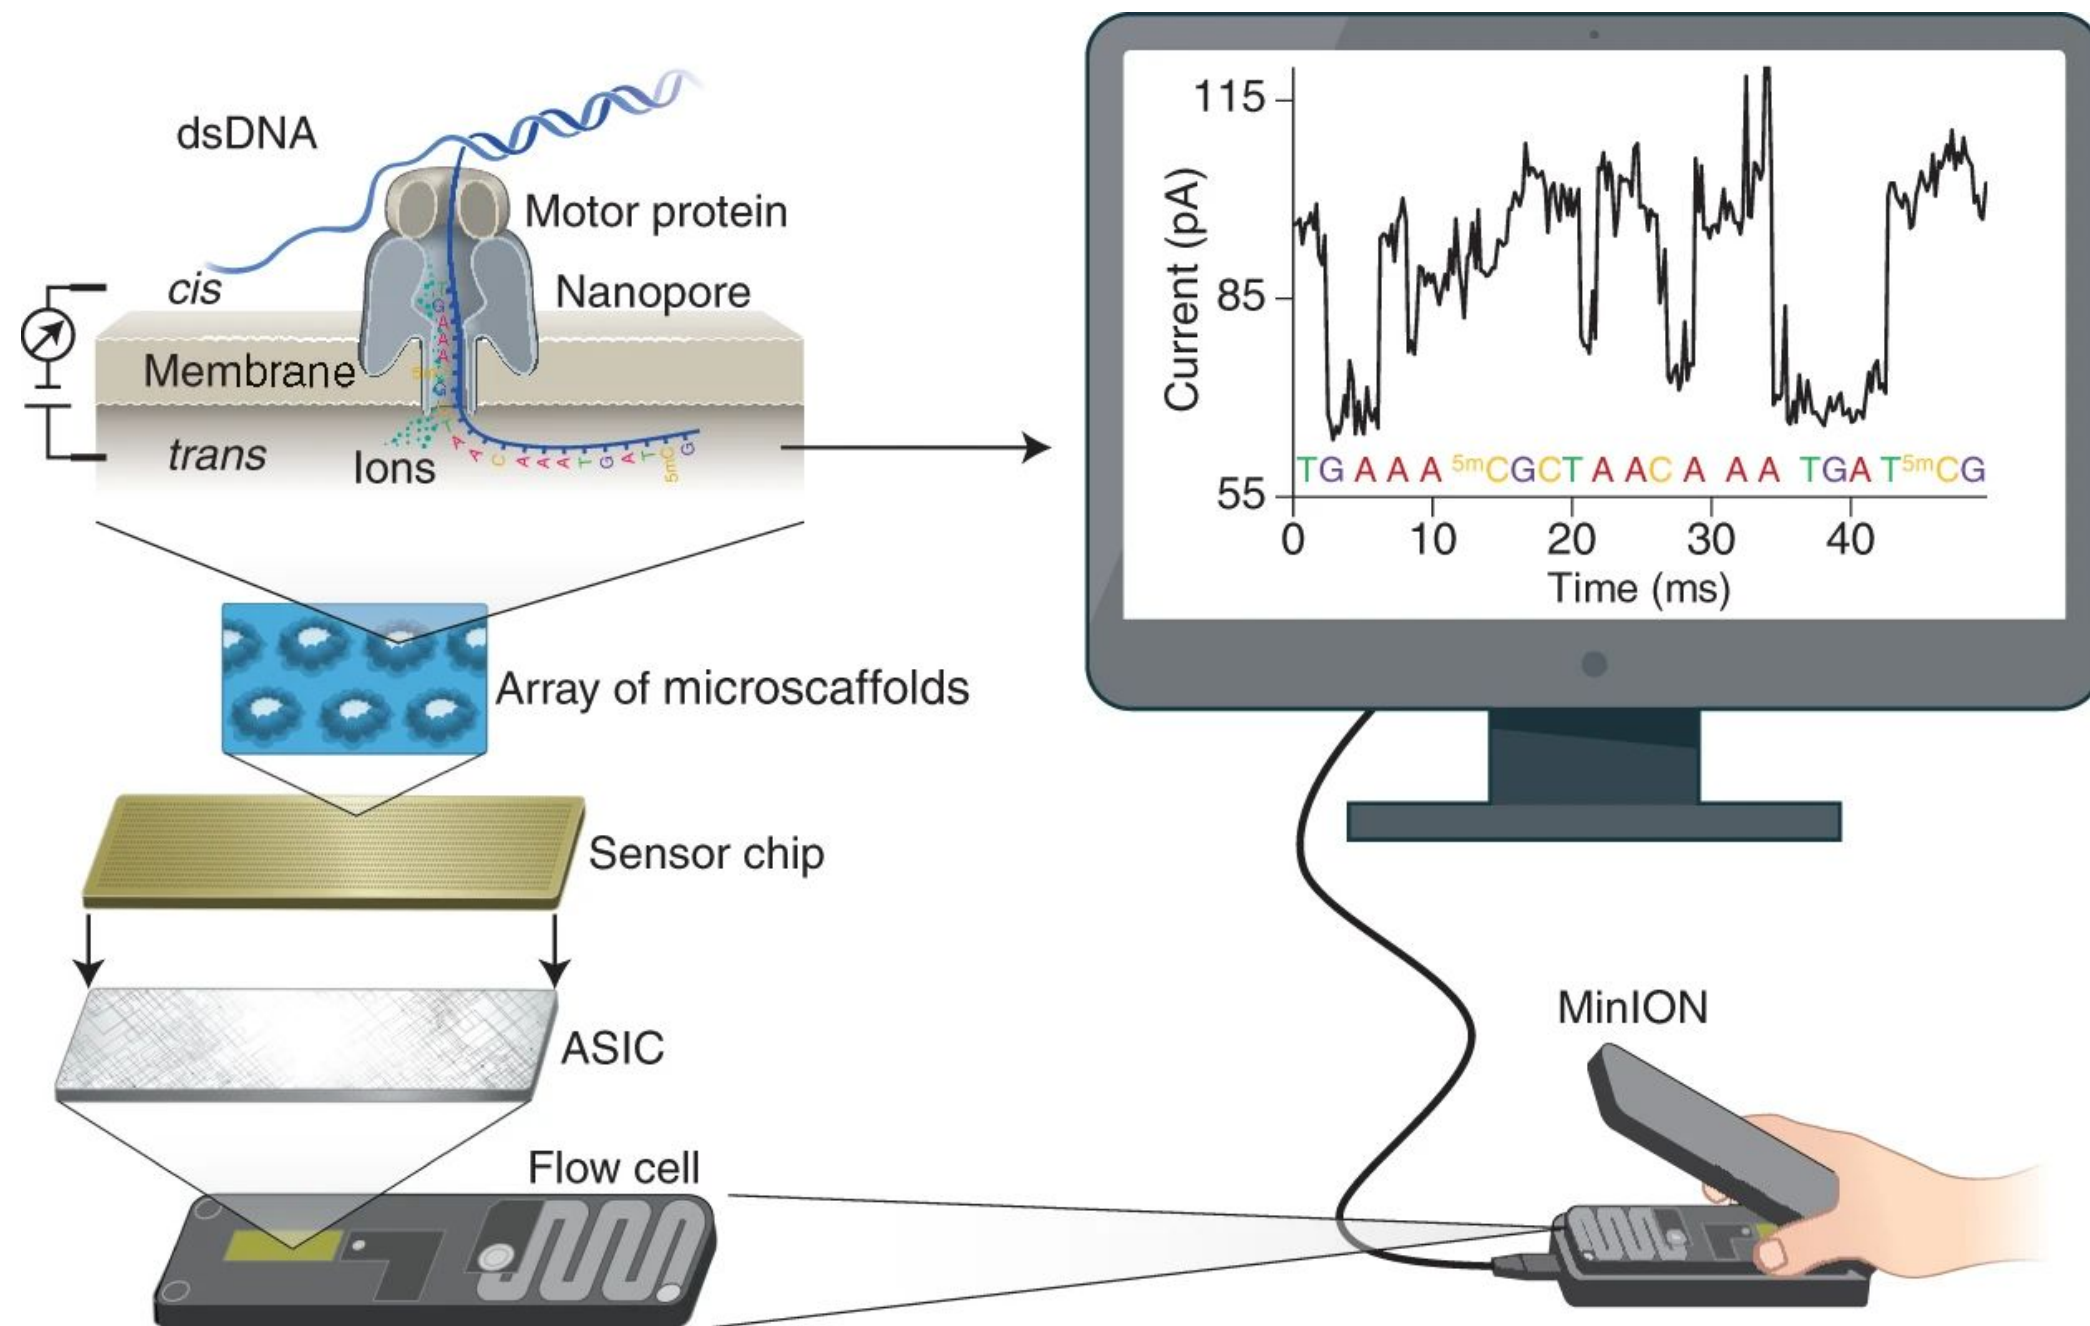
PacBio
Oxford Nanopore

PacBio Sequencing



<https://www.pacb.com/wp-content/uploads/SMRT-Sequencing-Brochure-Delivering-highly-accurate-long-reads-to-drive-discovery-in-life-science.pdf>

Nanopore Sequencing



<https://www.nature.com/articles/s41587-021-01108-x/figures/1>

Sequencing Technology

First generation



Sanger sequencing
Maxam and Gilbert
Sanger chain termination

Infer nucleotide identity using dNTPs,
then visualize with electrophoresis

500–1,000 bp fragments

Second generation (next generation sequencing)



454, Solexa,
Ion Torrent,
Illumina

High throughput from the
parallelization of sequencing reactions

~50–500 bp fragments

Third generation

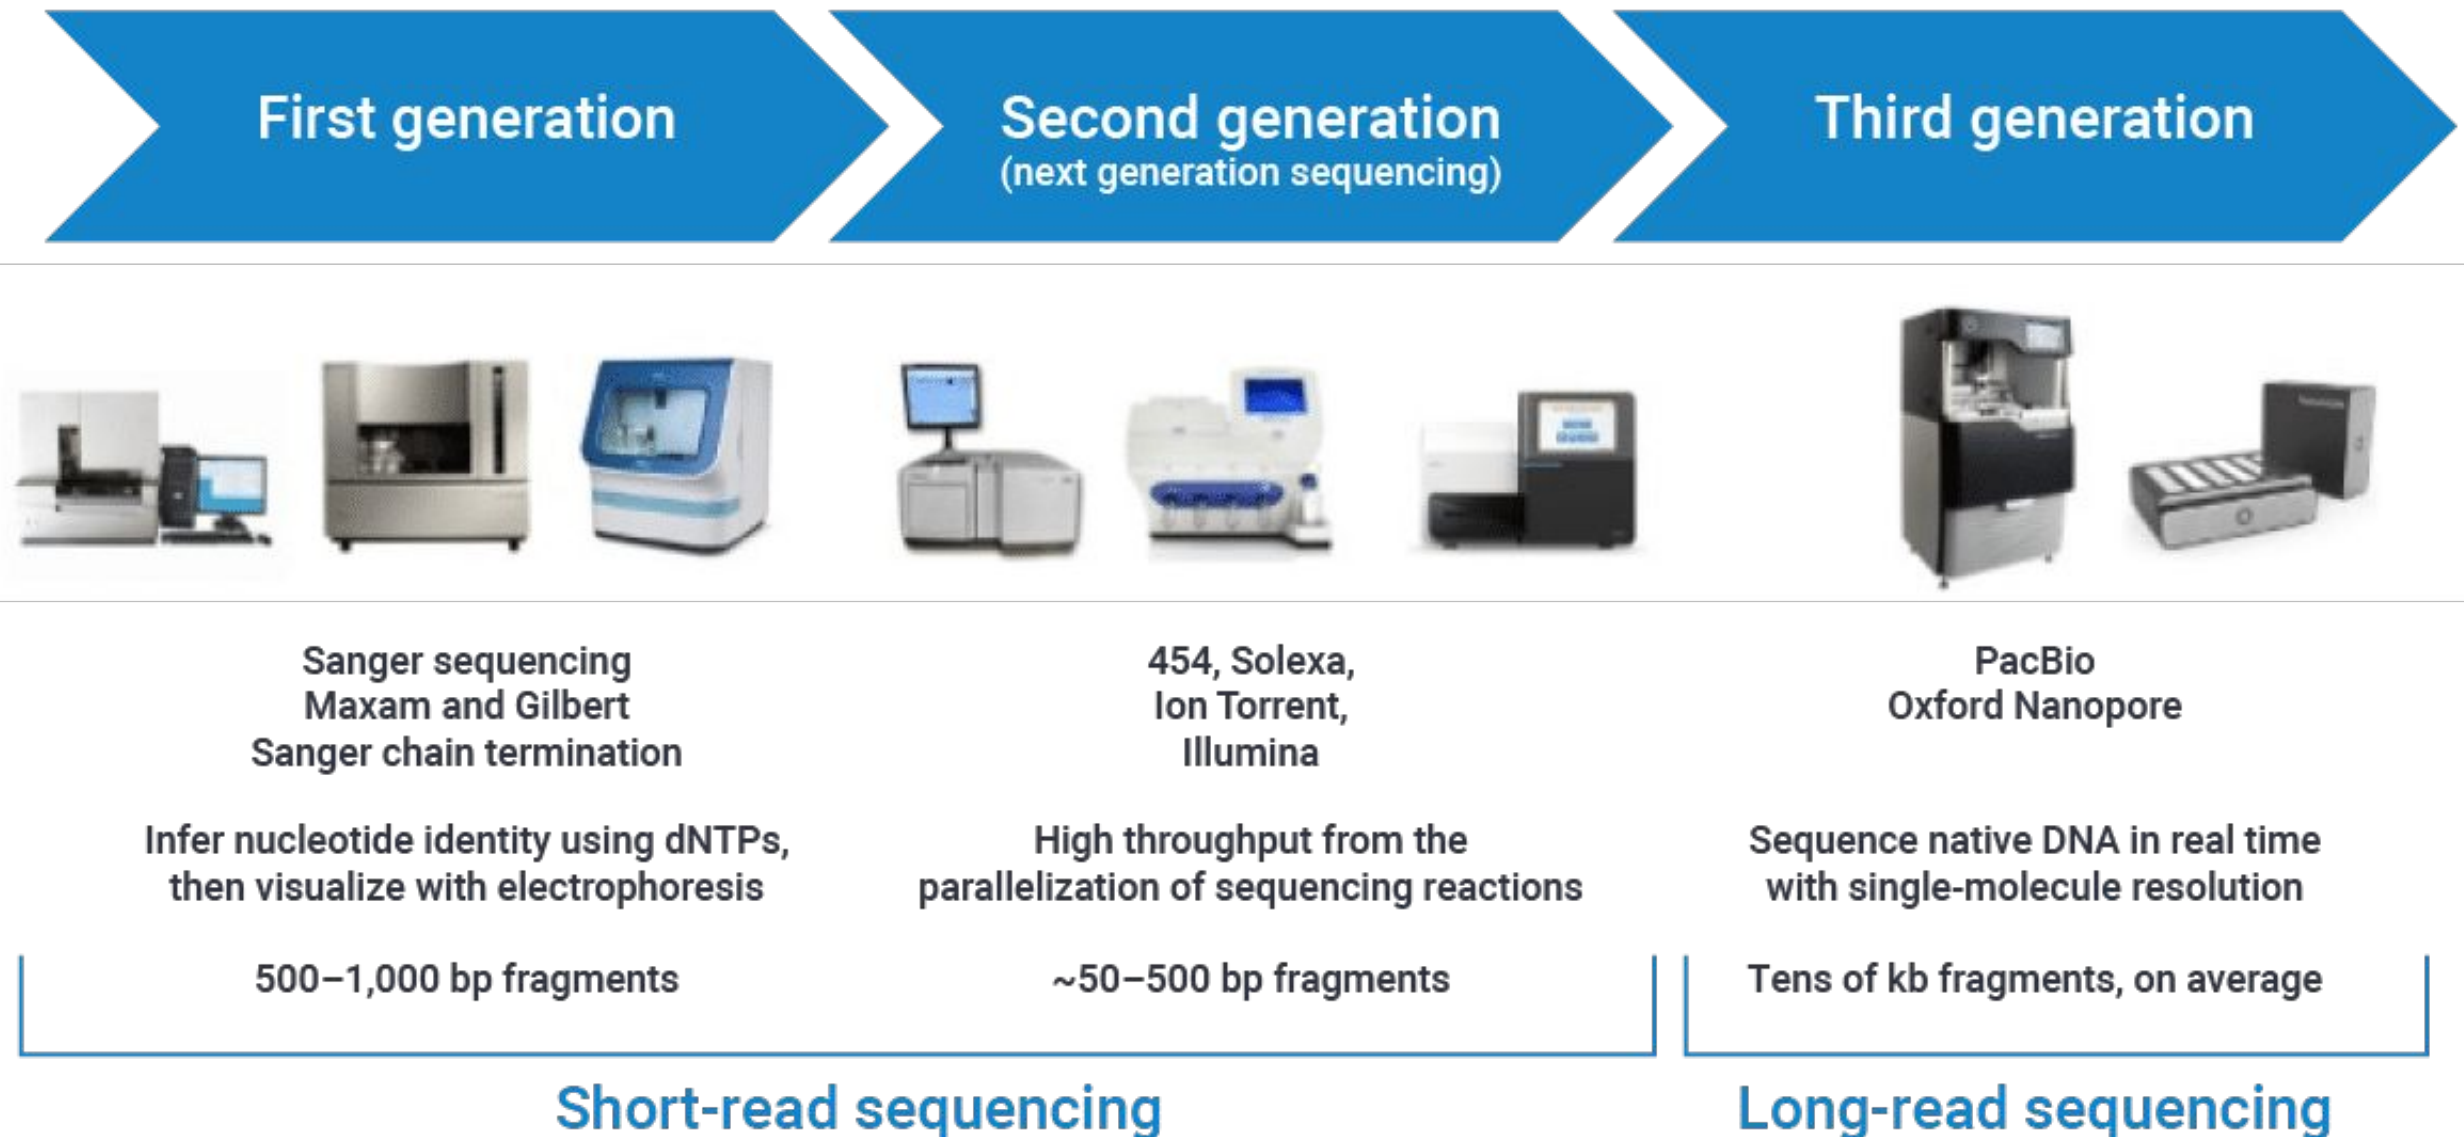


PacBio
Oxford Nanopore

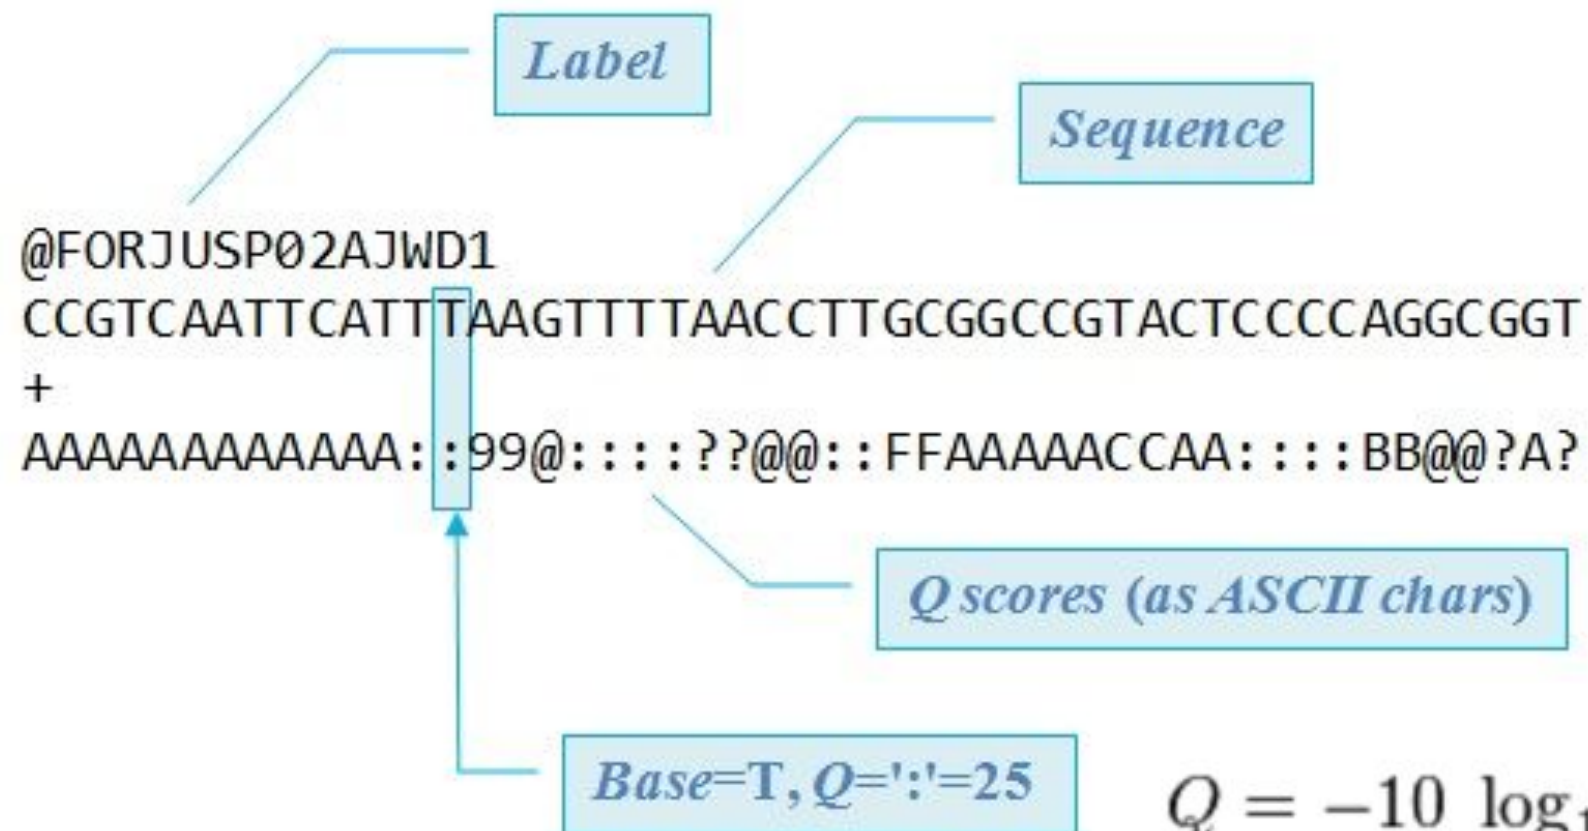
Sequence native DNA in real time
with single-molecule resolution

Tens of kb fragments, on average

Sequencing Technology



Capturing measurement error: FASTQ



Quality value Q is an integer representation of the probability p that a corresponding base call is incorrect

$$Q = -10 \log_{10} P \quad \longrightarrow \quad P = 10^{-\frac{Q}{10}}$$

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1000	99.9%
40	1 in 10000	99.99%
50	1 in 100000	99.999%

https://www.drive5.com/usearch/manual/fastq_files.html

<https://learn.gencore.bio.nyu.edu/ngs-file-formats/quality-scores/>


Assembly

Reads



+

Reference genome






Input DNA



How do we assemble puzzle without the benefit of knowing what the finished product should look like?

(That's what the Human Genome Project had to do!)

De novo shotgun assembly



Assembly

Whole-genome “shotgun” sequencing first copies the input DNA:

Input: GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT


Copy: GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTTT

Then fragments it:

Fragment: GGCGTCTA TATCTCGG CTCTAGGCCCTC ATTTTTTT
GGC GTCTATAT CTCGGCTCTAGGCCCTCA TTTTTT
GGCGTC TATATCT CGGCTCTAGGCCCT CATTTTTTT
GGCGTCTAT ATCTCGGCTCTAG GCCCTCA TTTTTT

“Shotgun” refers to the random fragmentation of the whole genome; like it was fired from a shotgun

Assembly




Assembly

CTAGGCCCTCAATTTT
GGCGTCTATATCT
CTCTAGGCCCTCAATTTT
TCTATATCTCGGCTCTAGG
GGCTCTAGGCCCTCATTTTT
CTCGGCTCTAGCCCCTCATT
TATCTCGACTCTAGGCCCTCA
GGCGTCGATATCT
TATCTCGACTCTAGGCC
GGCGTCTATATCTCG



From
these

Reconstruct this



??

Coverage

CTAGGCCCTCAATTTT
CTCTAGGCCCTCAATTTT
GGCTCTAGGCCCTCAATTTT
CTCGGCTCTAGGCCCTCAATTT
TATCTCGACTCTAGGCCCTCA
TATCTCGACTCTAGGCC
TCTATATCTCGGCTCTAGG
GGCGTCTATATCTCG
GGCGTCGATATCT
GGCGTCTATATCT
GGCGTCTATATCTCGGCTCTAGGCCCTCAATTTT

Coverage = 5

Coverage

CTAGGCCCTCAATTTT
CTCTAGGCCCTCAATTTT
GGCTCTAGGCCCTCAATTTT
CTCGGCTCTAGGCCCTCAATTT
TATCTCGACTCTAGGCCCTCA
TATCTCGACTCTAGGCC
TCTATATCTCGGCTCTAGG
GGCGTCTATATCTCG
GGCGTCGATATCT
GGCGTCTATATCT
GGCGTCTATATCTCGGCTCTAGGCCCTCAATTTT

Coverage = 5

CTAGGCCCTCAATTTT
CTCTAGGCCCTCAATTTT
GGCTCTAGGCCCTCAATTTT
CTCGGCTCTAGGCCCTCAATTT
TATCTCGACTCTAGGCCCTCA
TATCTCGACTCTAGGCC
TCTATATCTCGGCTCTAGG
GGCGTCTATATCTCG
GGCGTCGATATCT
GGCGTCTATATCT
GGCGTCTATATCTCGGCTCTAGGCCCTCAATTTT

177 bases

35 bases

Average coverage = $177 / 35 \approx 5$ -fold

TCTATATCTCGGCTCTAGG

TATCTCGACTCTAGGCC

TCTATATCTCGGCTCTAGG

|||||

TATCTCGACTCTAGGCC

First law of assembly

If a suffix of read A is similar to a prefix of read B...

```
TCTATATCTCGGCTCTAGG
      ||||| |||||
TATCTCGACTCTAGGCC
```

...then A and B might *overlap* in the genome

```
TCTATATCTCGGCTCTAGG
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTT
TATCTCGACTCTAGGCC
```

TCTATATCTCGGCTCTAGG
 |||||
 TATCTCGA CTCTAGGCC

Why the differences?

1. Sequencing errors
2. Ploidy: e.g. humans have 2 copies of each chromosome, and copies can differ



Second law of assembly

More coverage leads to more and longer overlaps

CTAGGCCCTCAATTTT
CTCGGCTCTAGCCCCTCATTTT
TCTATATCTCGGCTCTAGG less coverage
GGCGTCGATATCT
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTTT
CTAGGCCCTCAATTTT
GGCTCTAGGCCCTCATTTTTT
CTCGGCTCTAGCCCCTCATTTT
TATCTCGACTCTAGGCCCTCA
TCTATATCTCGGCTCTAGG more coverage
GGCGTCTATATCTCG
GGCGTCTATATCT

TCTATATCTCGGCTCTAGG

|||||

TATCTCGACTCTAGGCC

TCTATATCTCGGCTCTAGG

||||| |||||


TATCTCGACTCTAGGCC

TATCTCGACTCTAGGCC

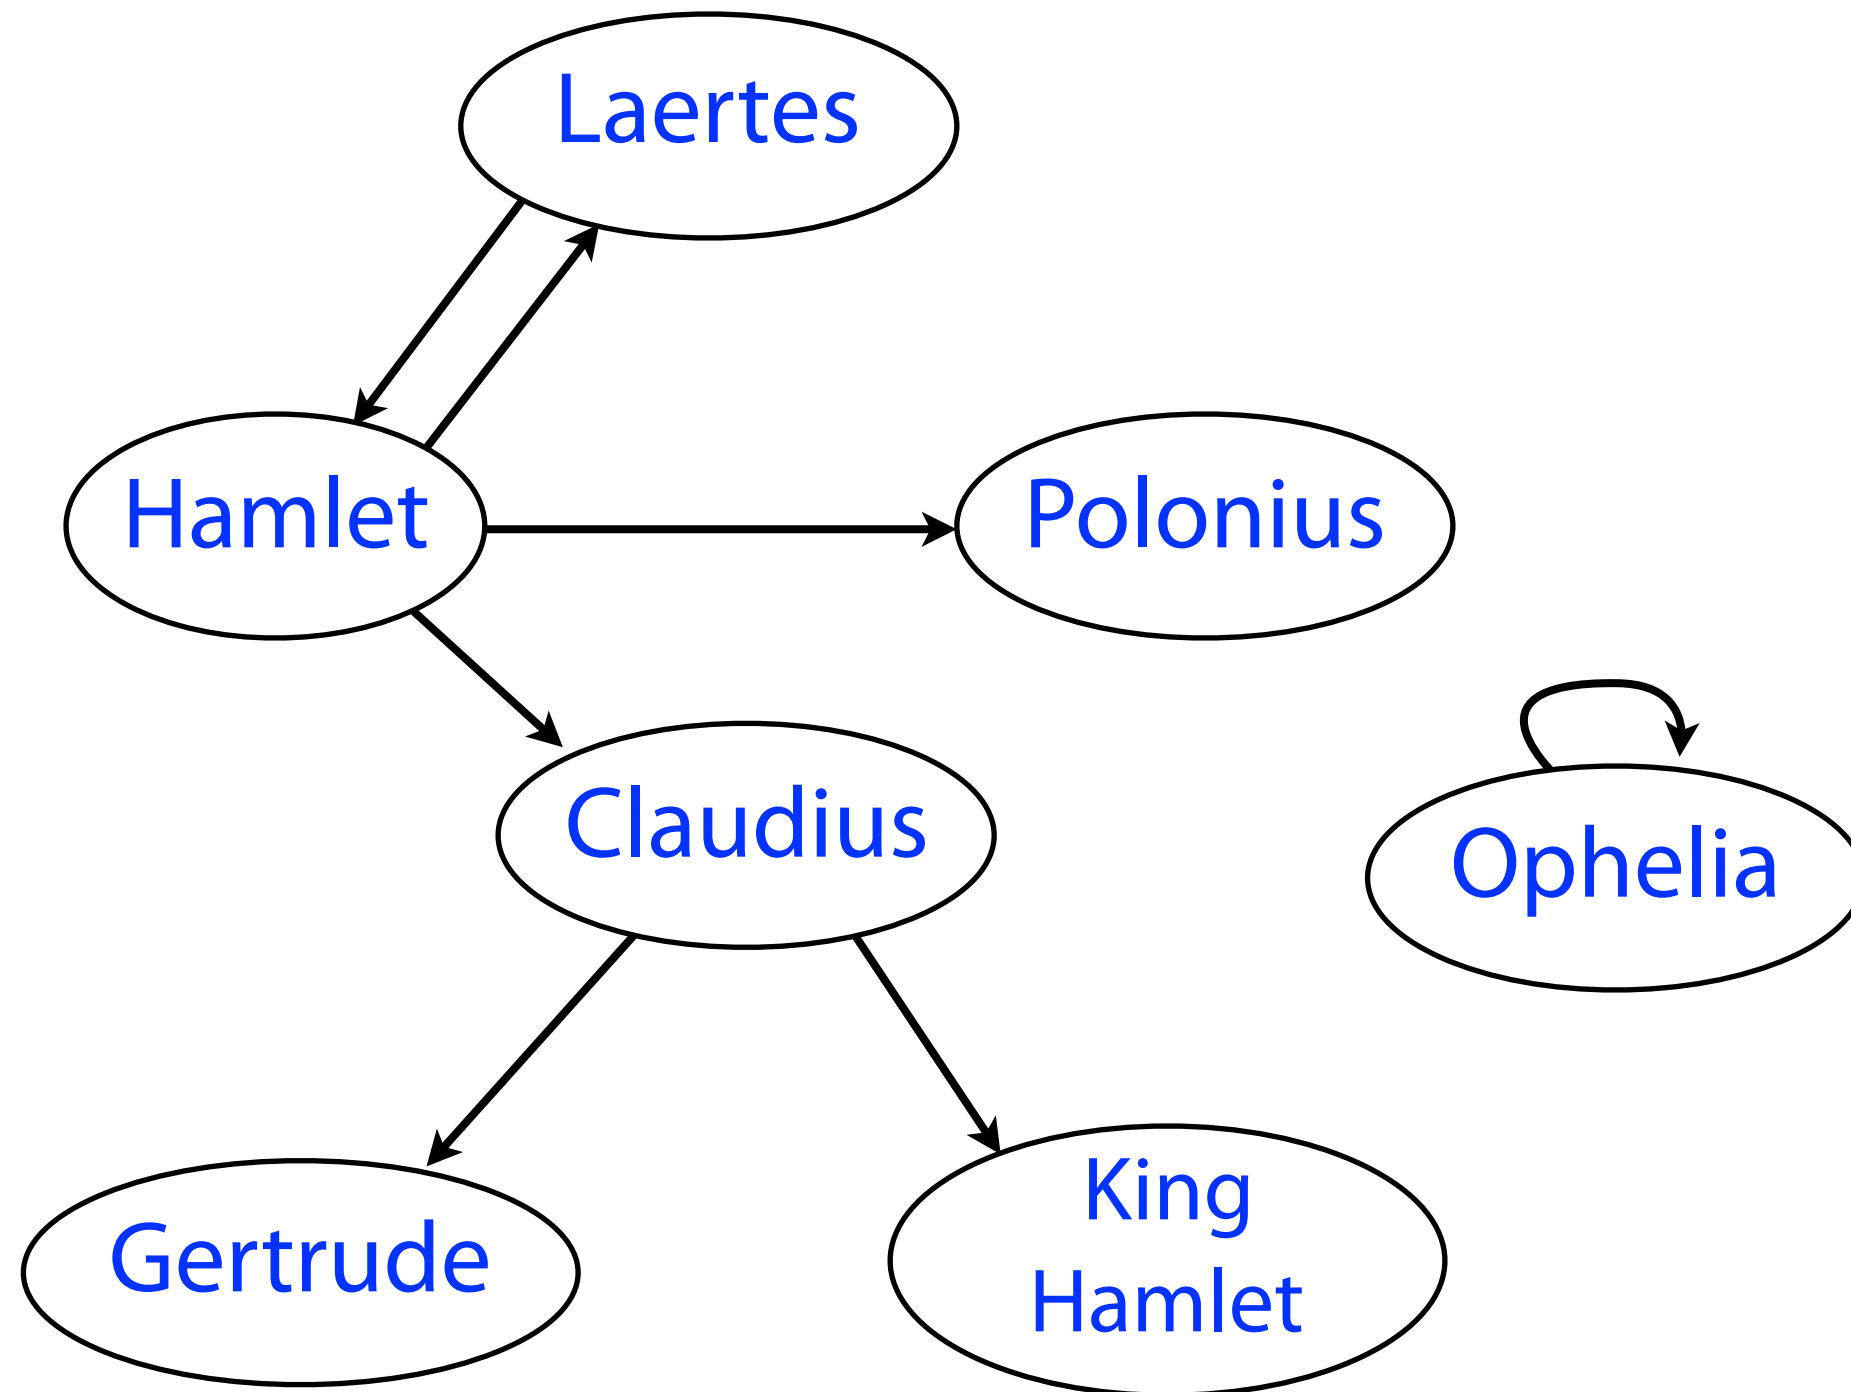
|||| | ||||| ||

CTCGGCTCTAGCCCTCAT

Directed graph



Directed graph



Overlap graph

Each node is a read

CTCGGGCTCTAGCCCCCTCATTTT

Draw edge A \rightarrow B when suffix of A overlaps prefix of B

CTCGGGCTCTAGCCCCCTCATTTT

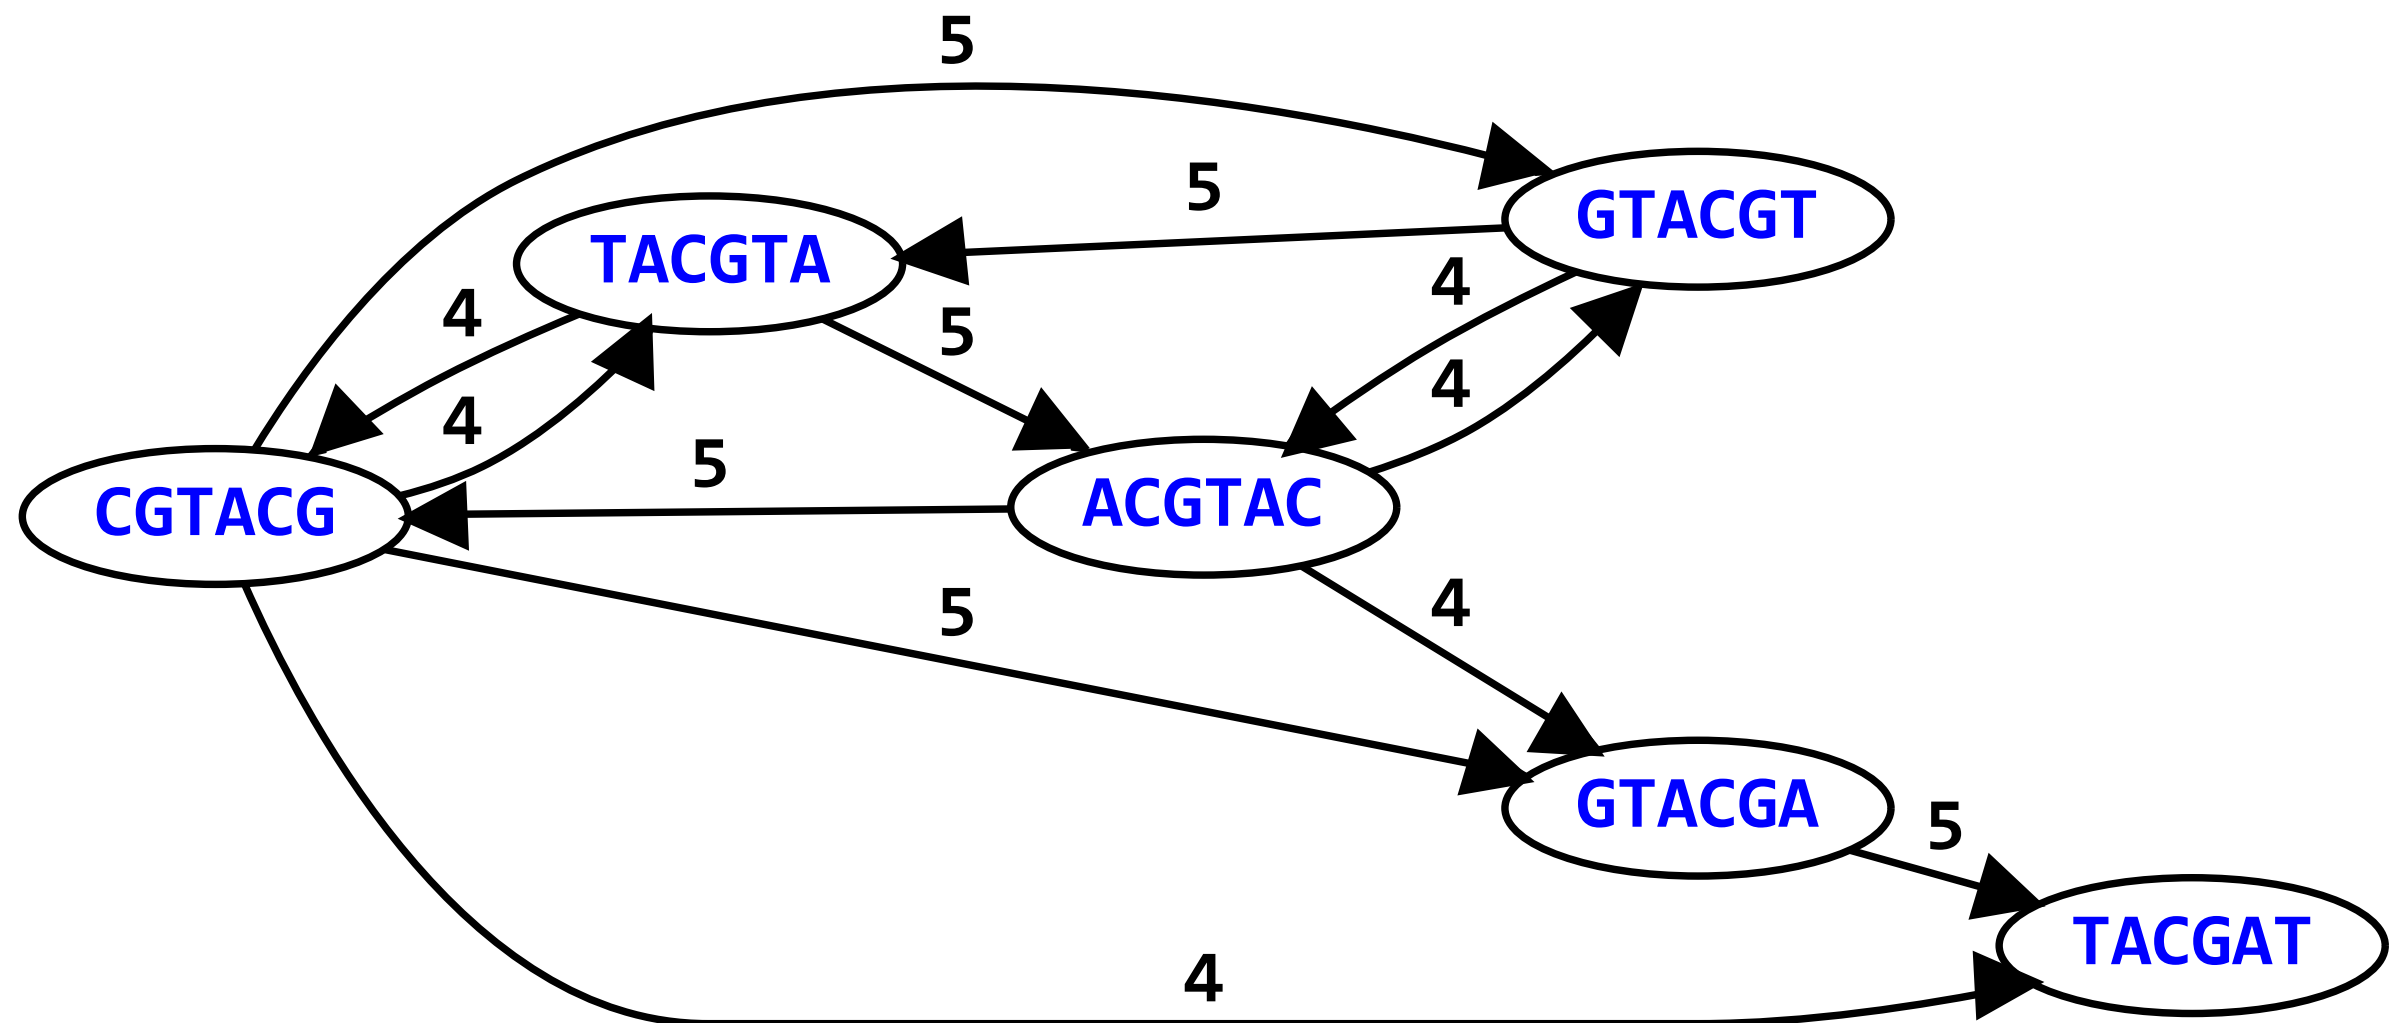


GGCTCTAGGCCCTCATTTTTT

Overlap graph

Nodes: all 6-mers from **GTACGTACGAT**

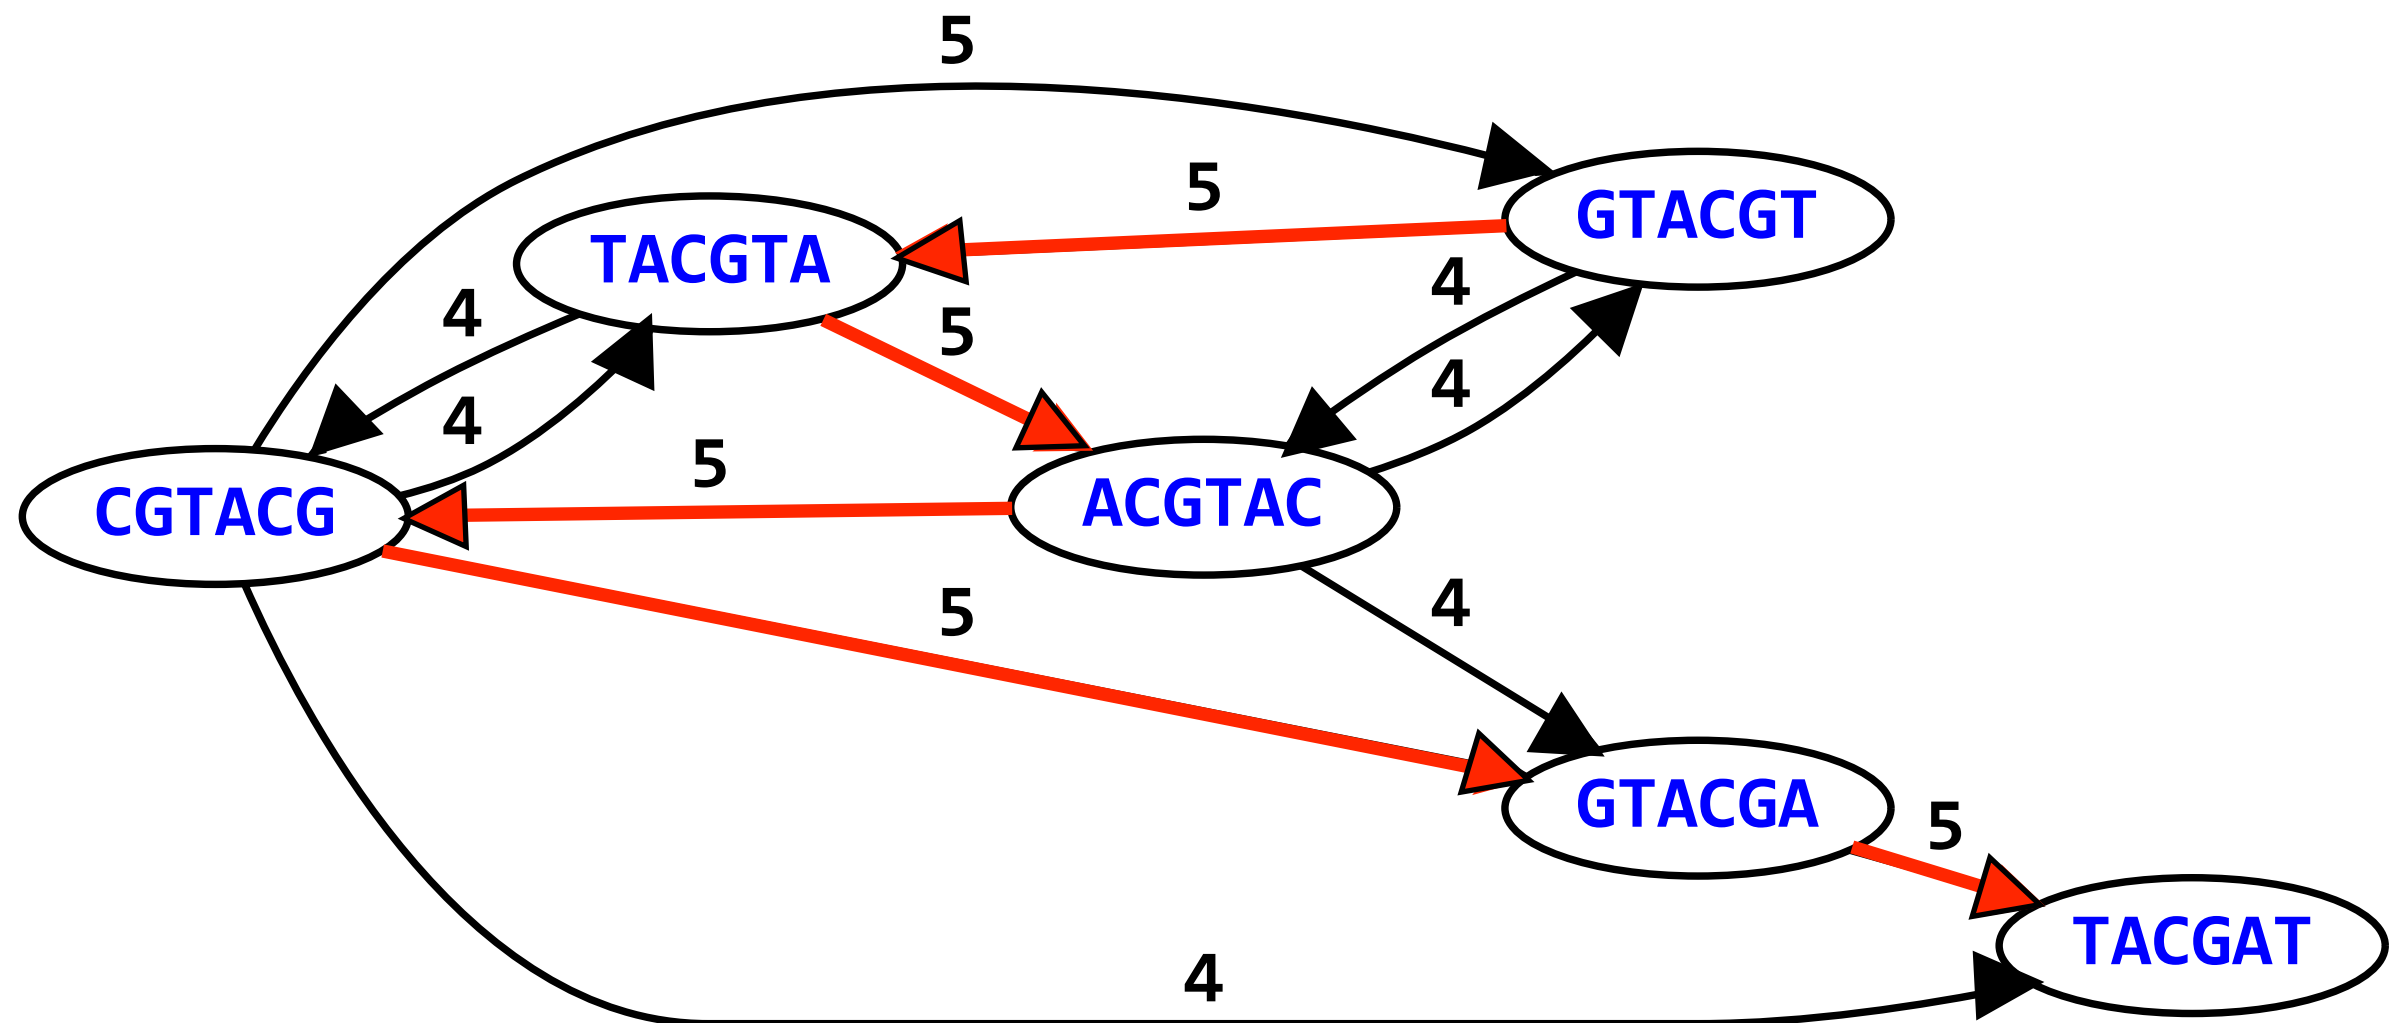
Edges: overlaps of length ≥ 4



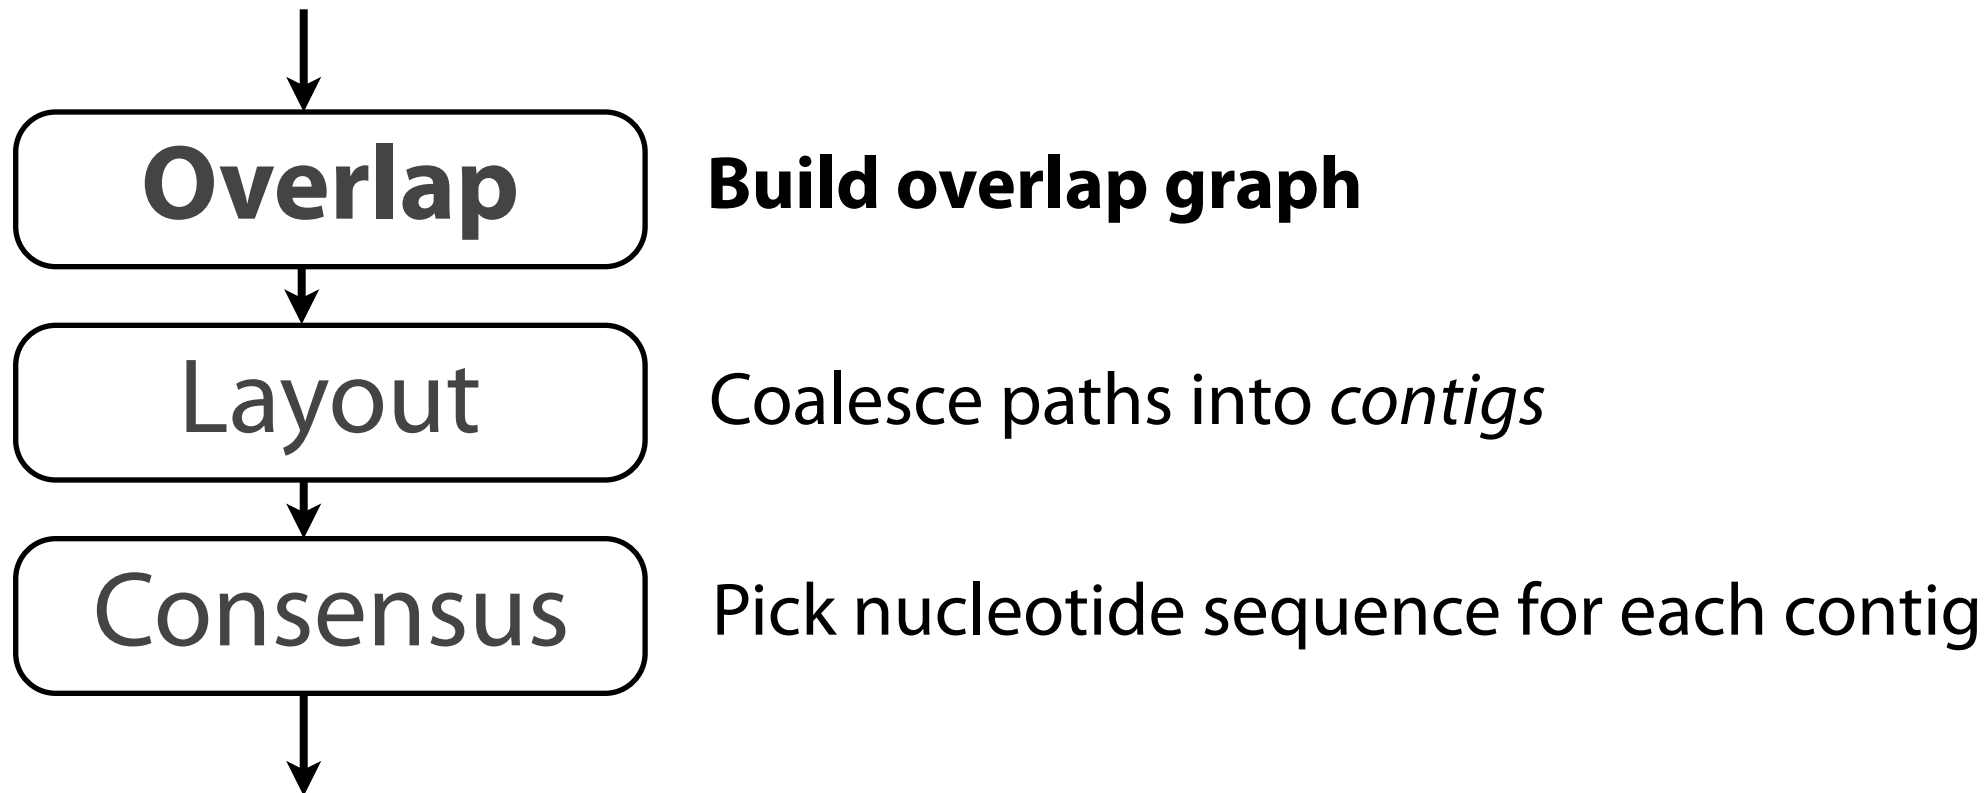
Overlap graph

Nodes: all 6-mers from **GTACGTACGAT**

Edges: overlaps of length ≥ 4



Overlap Layout Consensus



Finding overlaps

Overlap: Suffix of X of length $\geq l$ matches prefix of Y ; l is given

Naive: look in X for occurrences of Y 's length- l prefix. Extend matches to the right to confirm whether entire suffix of X matches.

Say $l = 3$

X : CTCTAGGCC

Y : TAGGCCCTC

Look for this in X

X : CTCTAGGCC

Y : TAGGCCCTC

Found it

Extend to right; confirm a length-6 prefix of Y matches a suffix of X

X : CTCTAGGCC

Y : TAGGCCCTC

See suffixPrefixMatch function in HW5 Q4 (Assembly Challenge)

Finding overlaps

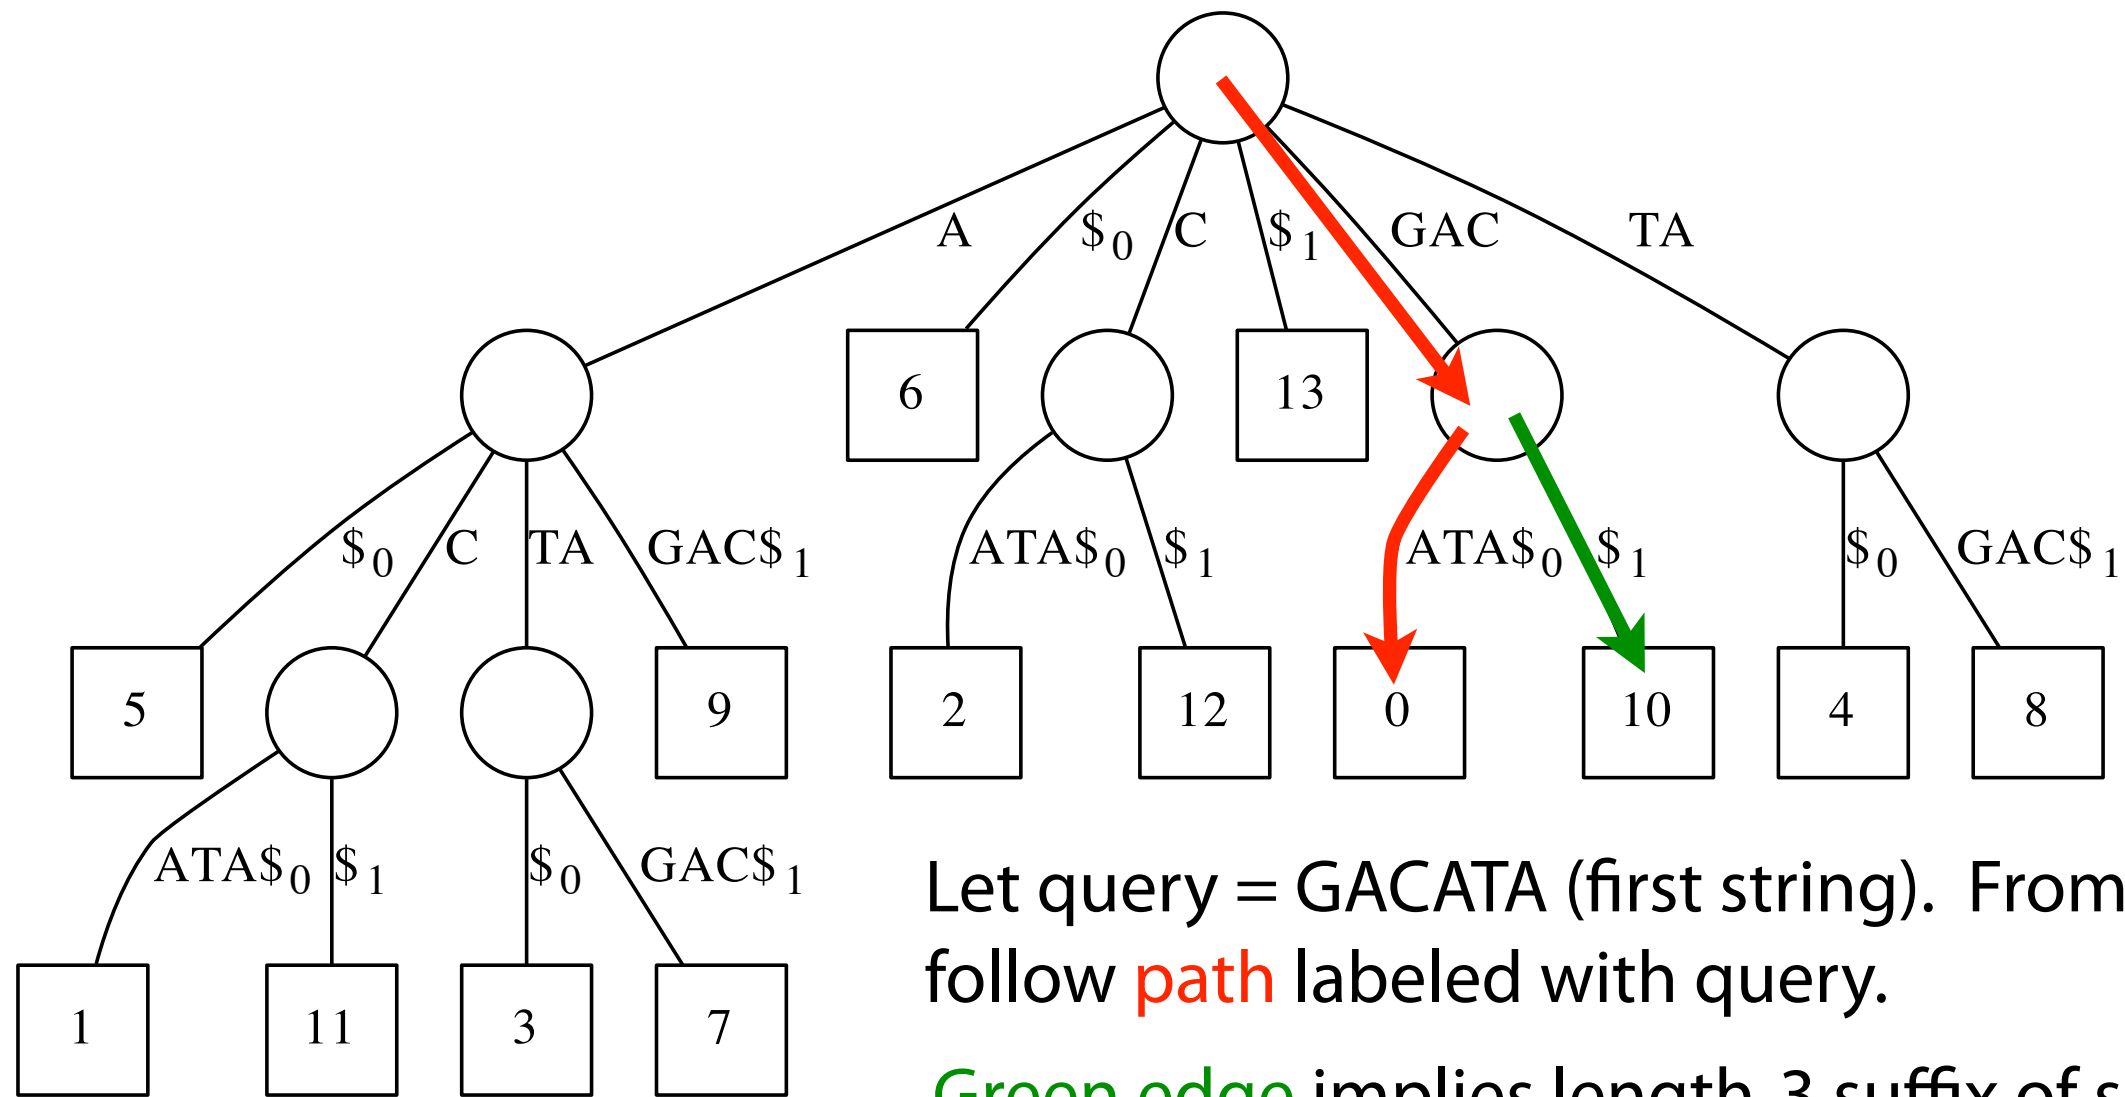
With suffix tree?

Given a collection of strings S , for each string x in S find all overlaps involving a prefix of x and a suffix of another string y

Finding overlaps with suffix tree

Generalized suffix tree for { "GACATA", "ATAGAC" }

GACATA\$₀ATAGAC\$₁



Let query = GACATA (first string). From root, follow **path** labeled with query.

Green edge implies length-3 suffix of second string equals length-3 prefix of query

ATAGAC

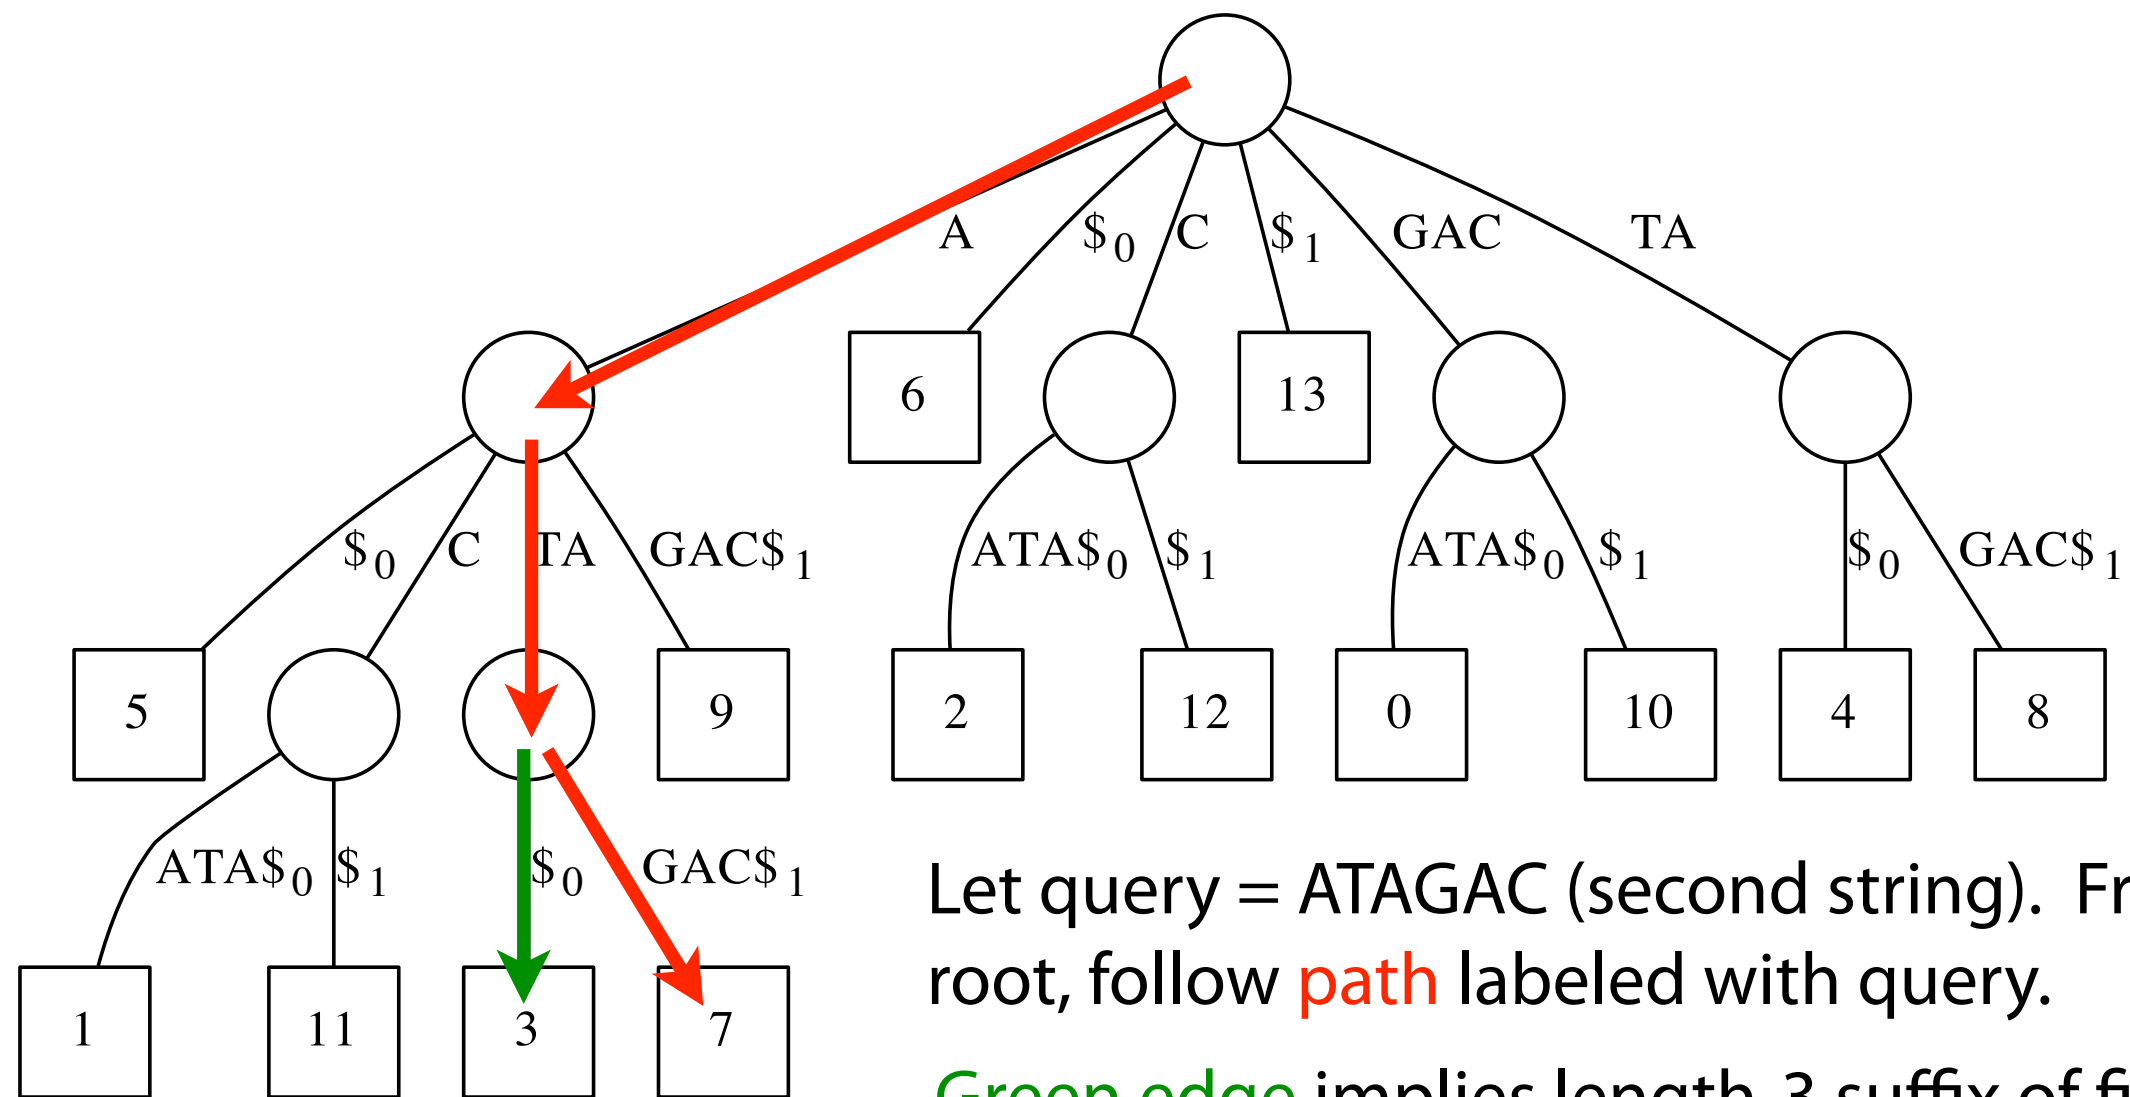
|||

GACATA

Finding overlaps with suffix tree

Generalized suffix tree for {"GACATA", "ATAGAC"}

GACATA\$₀ATAGAC\$₁



GACATA
|||
ATAGAC

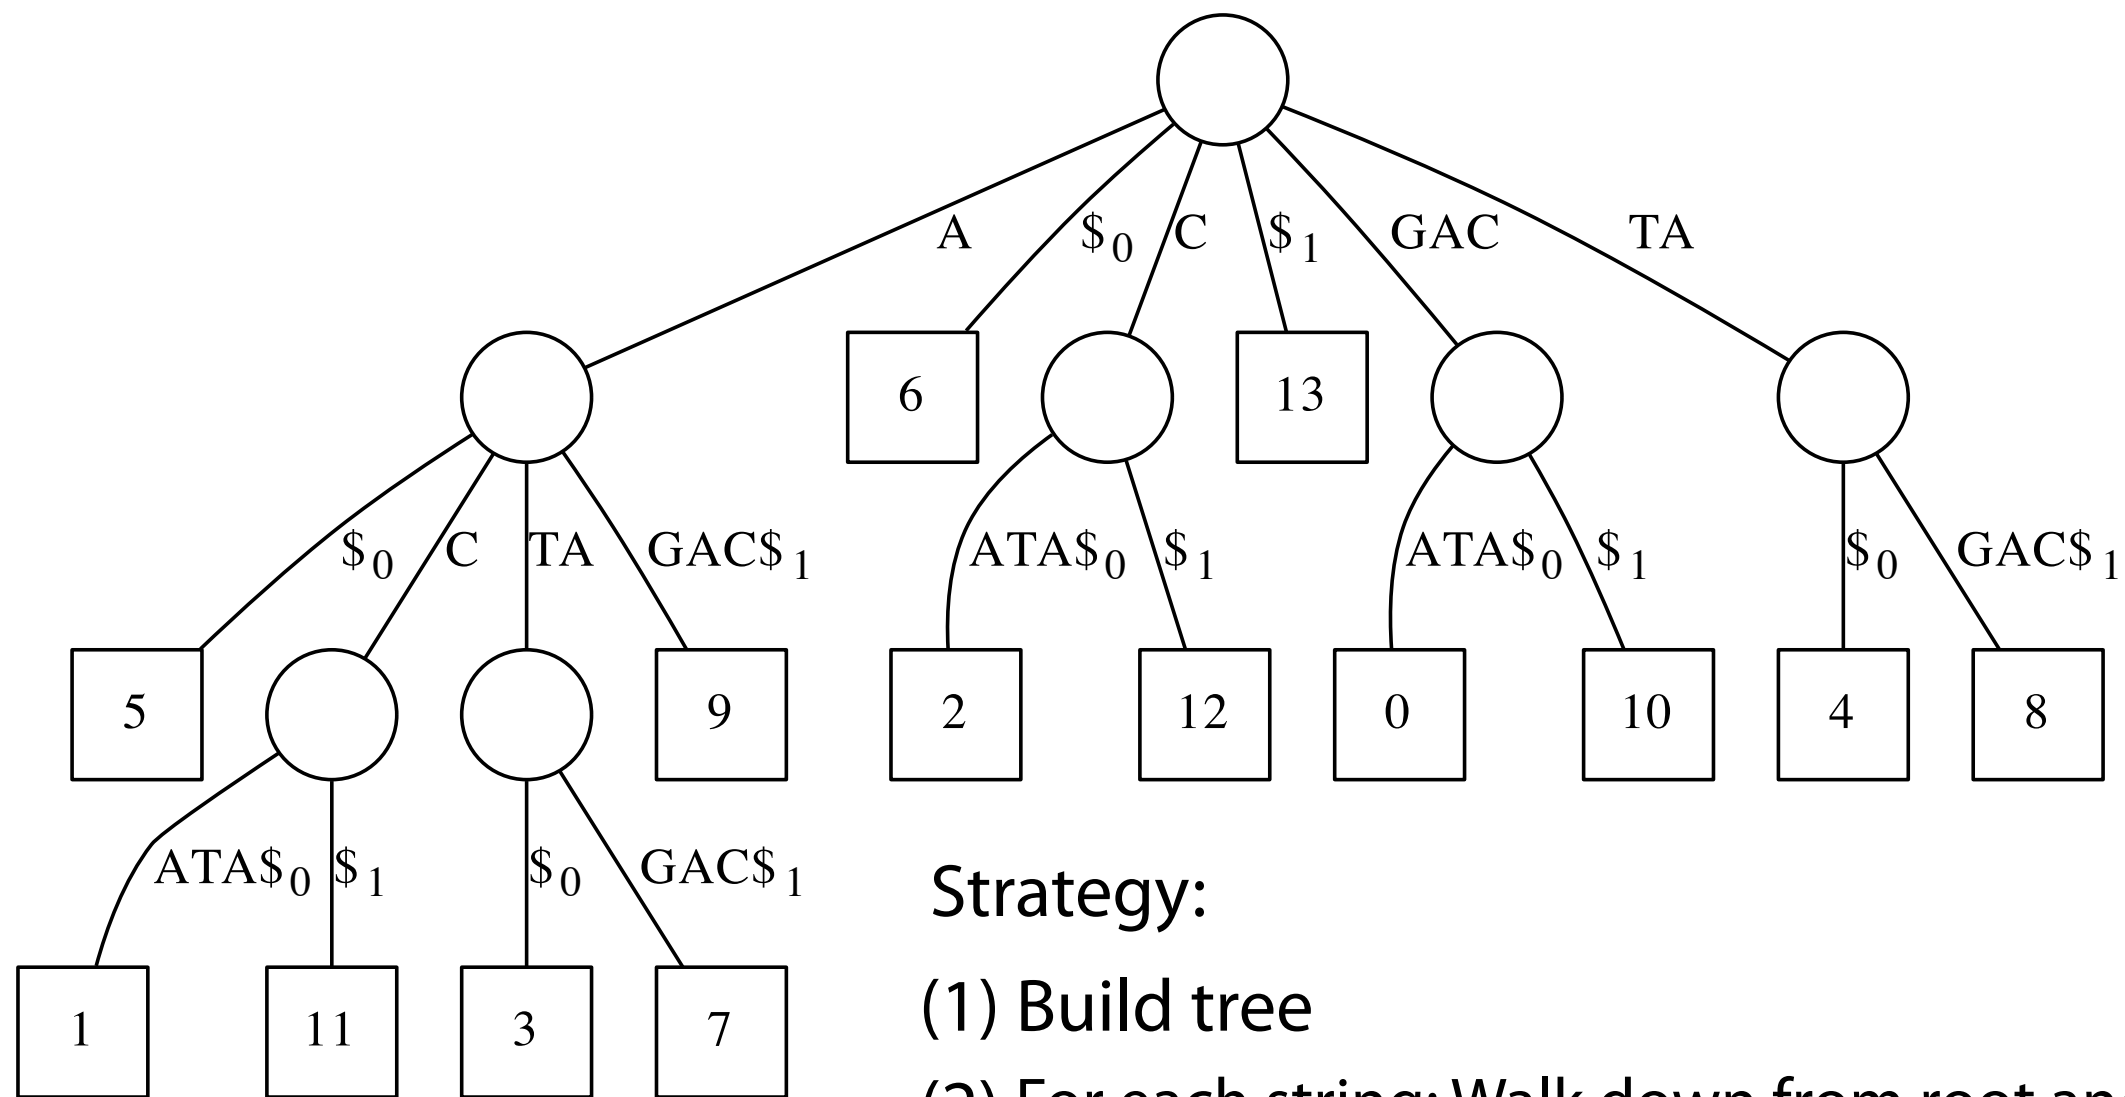
Let query = ATAGAC (second string). From root, follow **path** labeled with query.

Green edge implies length-3 suffix of first string equals length-3 prefix of query

Finding overlaps with suffix tree

Generalized suffix tree for {"GACATA", "ATAGAC"}

GACATA\$₀ATAGAC\$₁

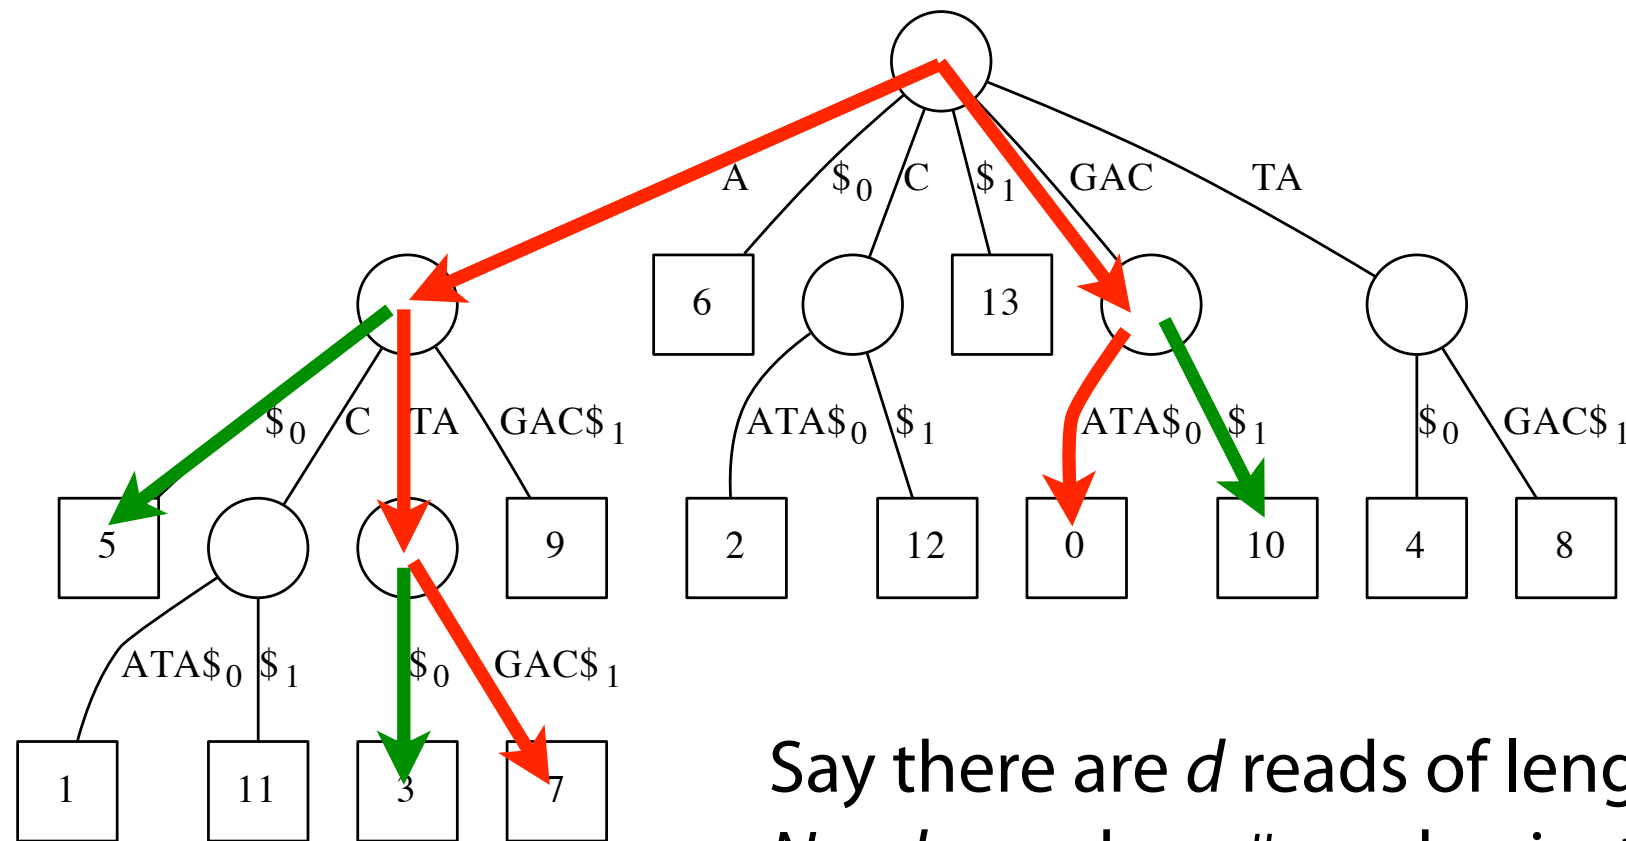


Strategy:

(1) Build tree

(2) For each string: Walk down from root and report any outgoing edge labeled with a separator. Each corresponds to a prefix/suffix match involving prefix of query string and suffix of string ending in the separator.

Finding overlaps with suffix tree



Say there are d reads of length n , total length $N = dn$, and $a = \#$ read pairs that overlap

Assume for given string pair we report only the longest suffix/prefix match

Time to build generalized suffix tree: $O(N)$

... to walk down red paths: $O(N)$

... to find & report overlaps (green): $O(a)$

Overall: $O(N + a)$

Finding overlaps

What about *approximate* suffix/prefix matches?

X: CTCGGCCCTAGG
 | | | | |
Y: GGCTCTAGGCC

Dynamic programming

Finding overlaps with dynamic programming

X: CTCGGCCCTAGG
 ||| ||||
Y: GGCTCTAGGCC

Use *global alignment* recurrence and score function

$$D[i, j] = \min \begin{cases} D[i-1, j] + s(x[i-1], -) \\ D[i, j-1] + s(-, y[j-1]) \\ D[i-1, j-1] + s(x[i-1], y[j-1]) \end{cases}$$

$s(a, b)$

	A	C	G	T	-
A	0	4	2	4	8
C	4	0	4	2	8
G	2	4	0	4	8
T	4	2	4	0	8
-	8	8	8	8	

How do we force it to find prefix / suffix matches?

Finding overlaps with dynamic programming

$$s(a, b)$$

	A	C	G	T	-
A	0	4	2	4	8
C	4	0	4	2	8
G	2	4	0	4	8
T	4	2	4	0	8
-	8	8	8	8	

How to initialize first row & column
so suffix of X aligns to prefix of Y ?

First column gets 0s
(any suffix of X is possible)

First row gets ∞ s
(must be a prefix of Y)

Backtrace from last row

Y

	-	G	G	C	T	C	T	A	G	G	C	C	C
-	0	∞	∞	∞	∞	∞	∞	∞	∞	∞	∞	∞	∞
C	0	4	12	20	28	36	44	52	60	68	76	84	92
T	0	4	8	14	22	30	38	46	54	62	70	78	86
C	0	4	8	8	16	24	32	40	48	56	64	72	80
G	0	0	4	12	16	16	24	26	30	36	44	52	60
G	0	0	0	8	16	16	24	26	30	36	44	52	60
C	0	4	4	0	8	16	18	26	30	34	36	44	52
C	0	4	8	4	0	8	16	22	30	34	34	36	44
C	0	4	8	8	6	0	10	18	26	34	34	34	36
T	0	4	8	10	8	8	0	10	18	26	34	36	36
A	0	2	6	12	14	12	10	0	10	18	26	34	40
G	0	0	2	10	16	18	16	10	0	10	18	26	34
G	0	0	0	6	14	20	22	18	10	0	10	18	26

X

X : CTCGGCCCTAGG

||| |||

Y : GGCTCTAGGCC

Finding overlaps with dynamic programming

Say there are d reads of length n , total length $N = dn$, and a is total number of pairs with an overlap

overlaps to try: $O(d^2)$

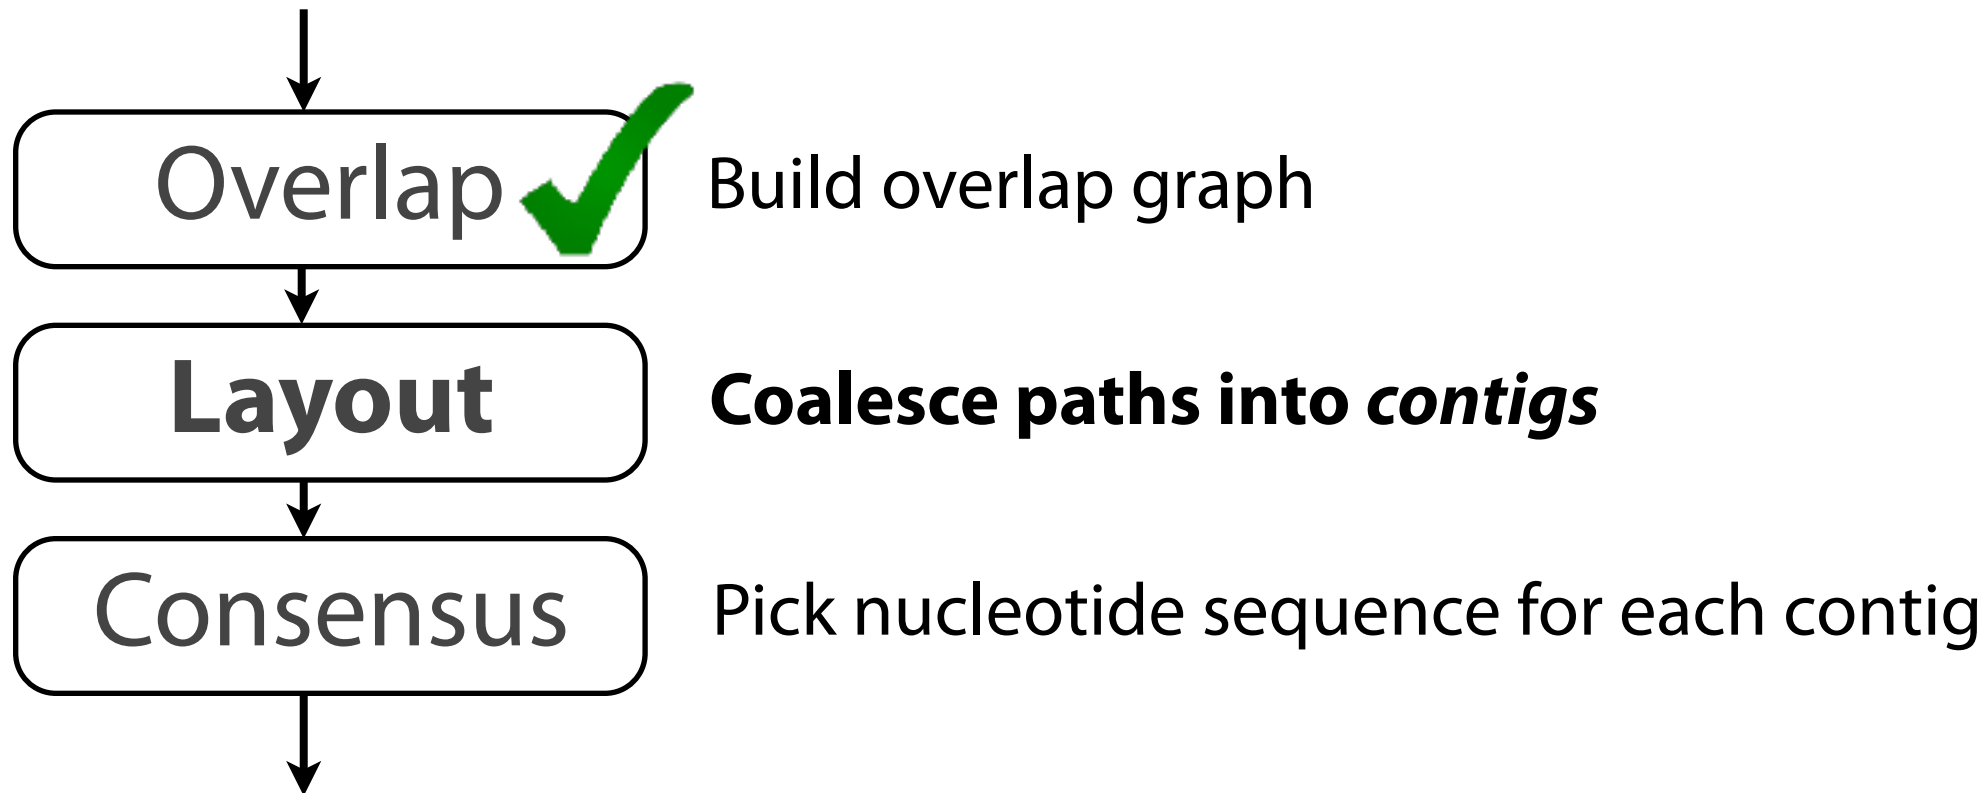
Size of each DP matrix: $O(n^2)$

Overall: $O(d^2n^2)$, or $O(N^2)$

Contrast $O(N^2)$ with suffix tree: $O(N + a)$, but where a is worst-case $O(d^2)$

Real-world overlappers mix the two; index filters out vast majority of non-overlapping pairs, dynamic programming used for remaining pairs

Overlap Layout Consensus



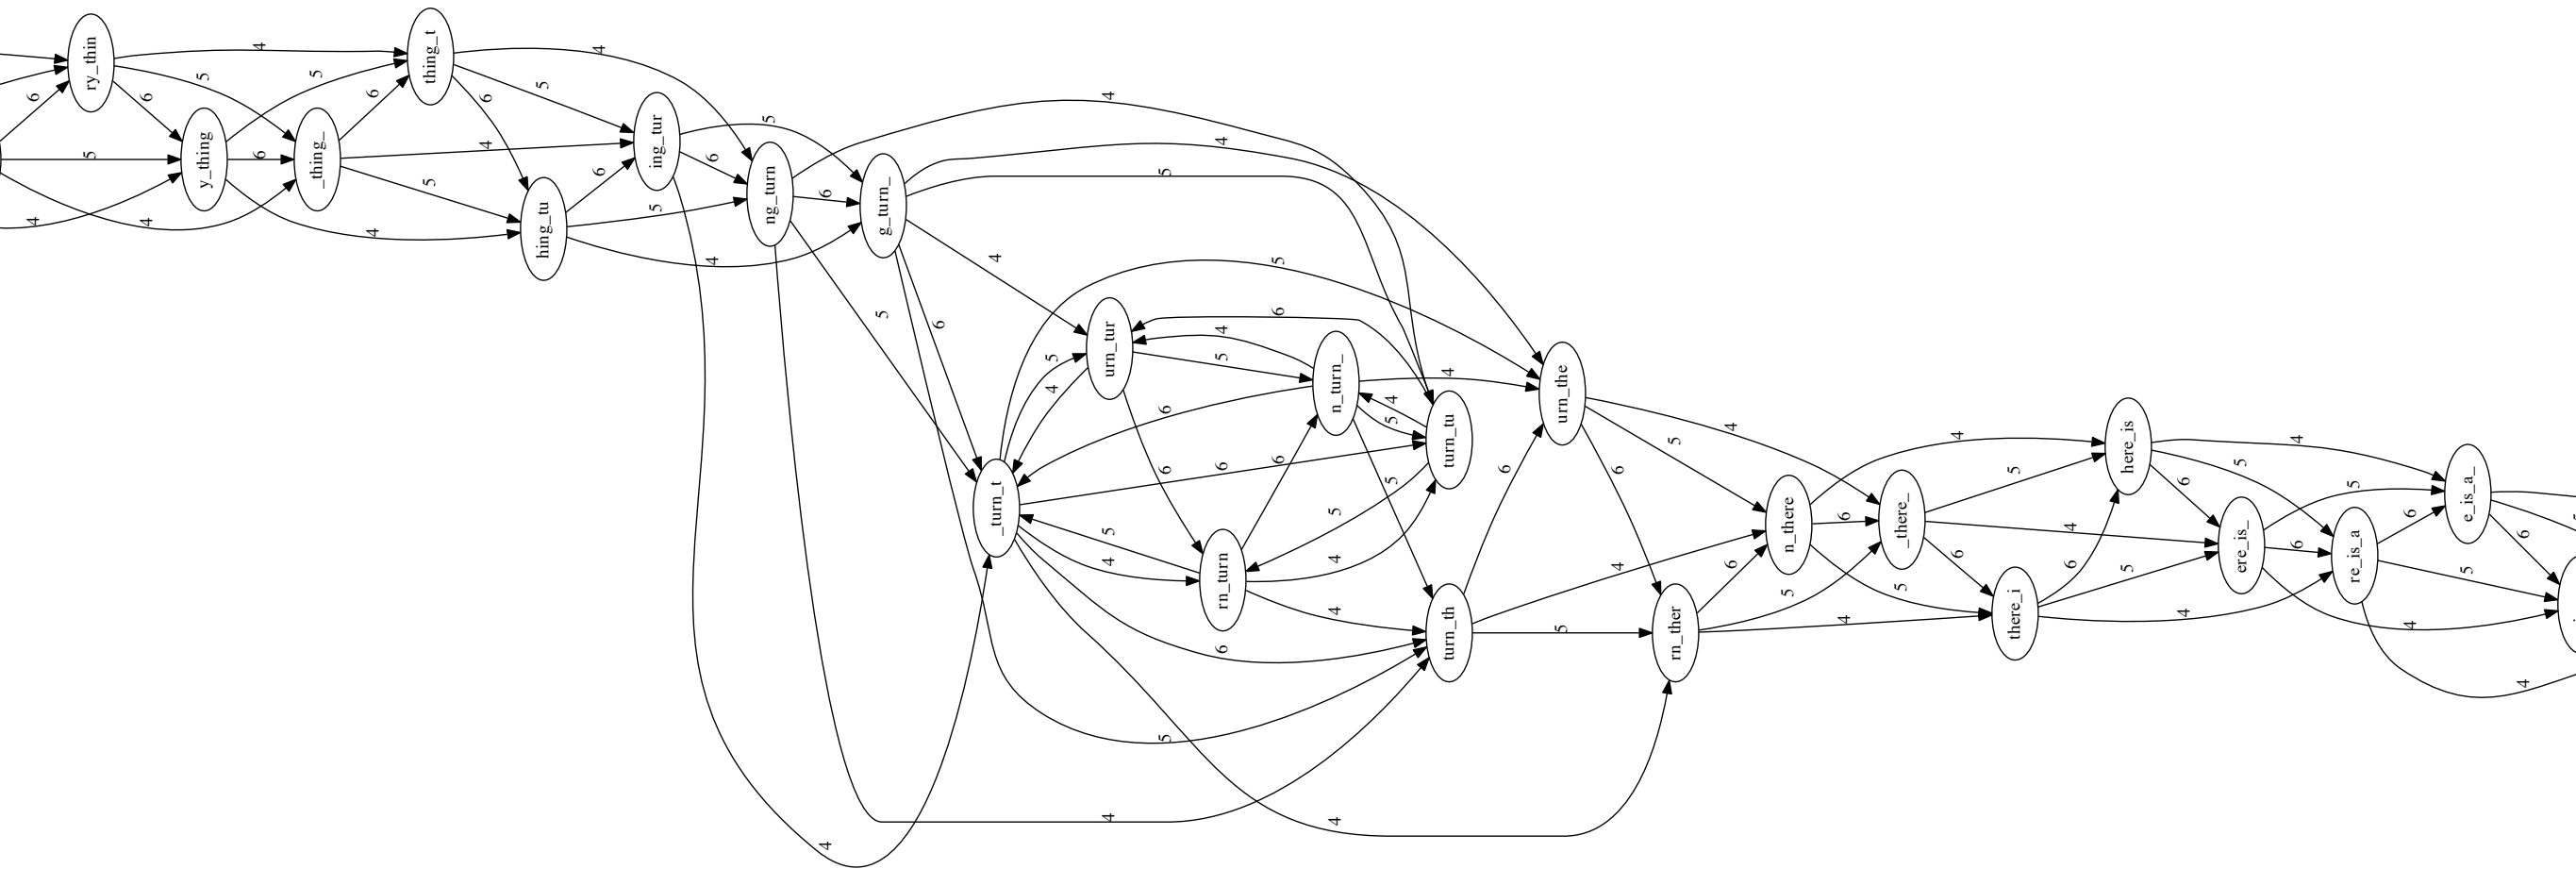
Layout

Overlap graph is big and messy. Contigs don't "pop out" at us.

Below: part of the overlap graph for

to_everything_turn_turn_turn_there_is_a_season

$l = 4, k = 7$




Layout

Anything redundant about this part of the overlap graph?


Some edges can be *inferred (transitively)* from other edges

E.g. **green** edge can be inferred from **blue**

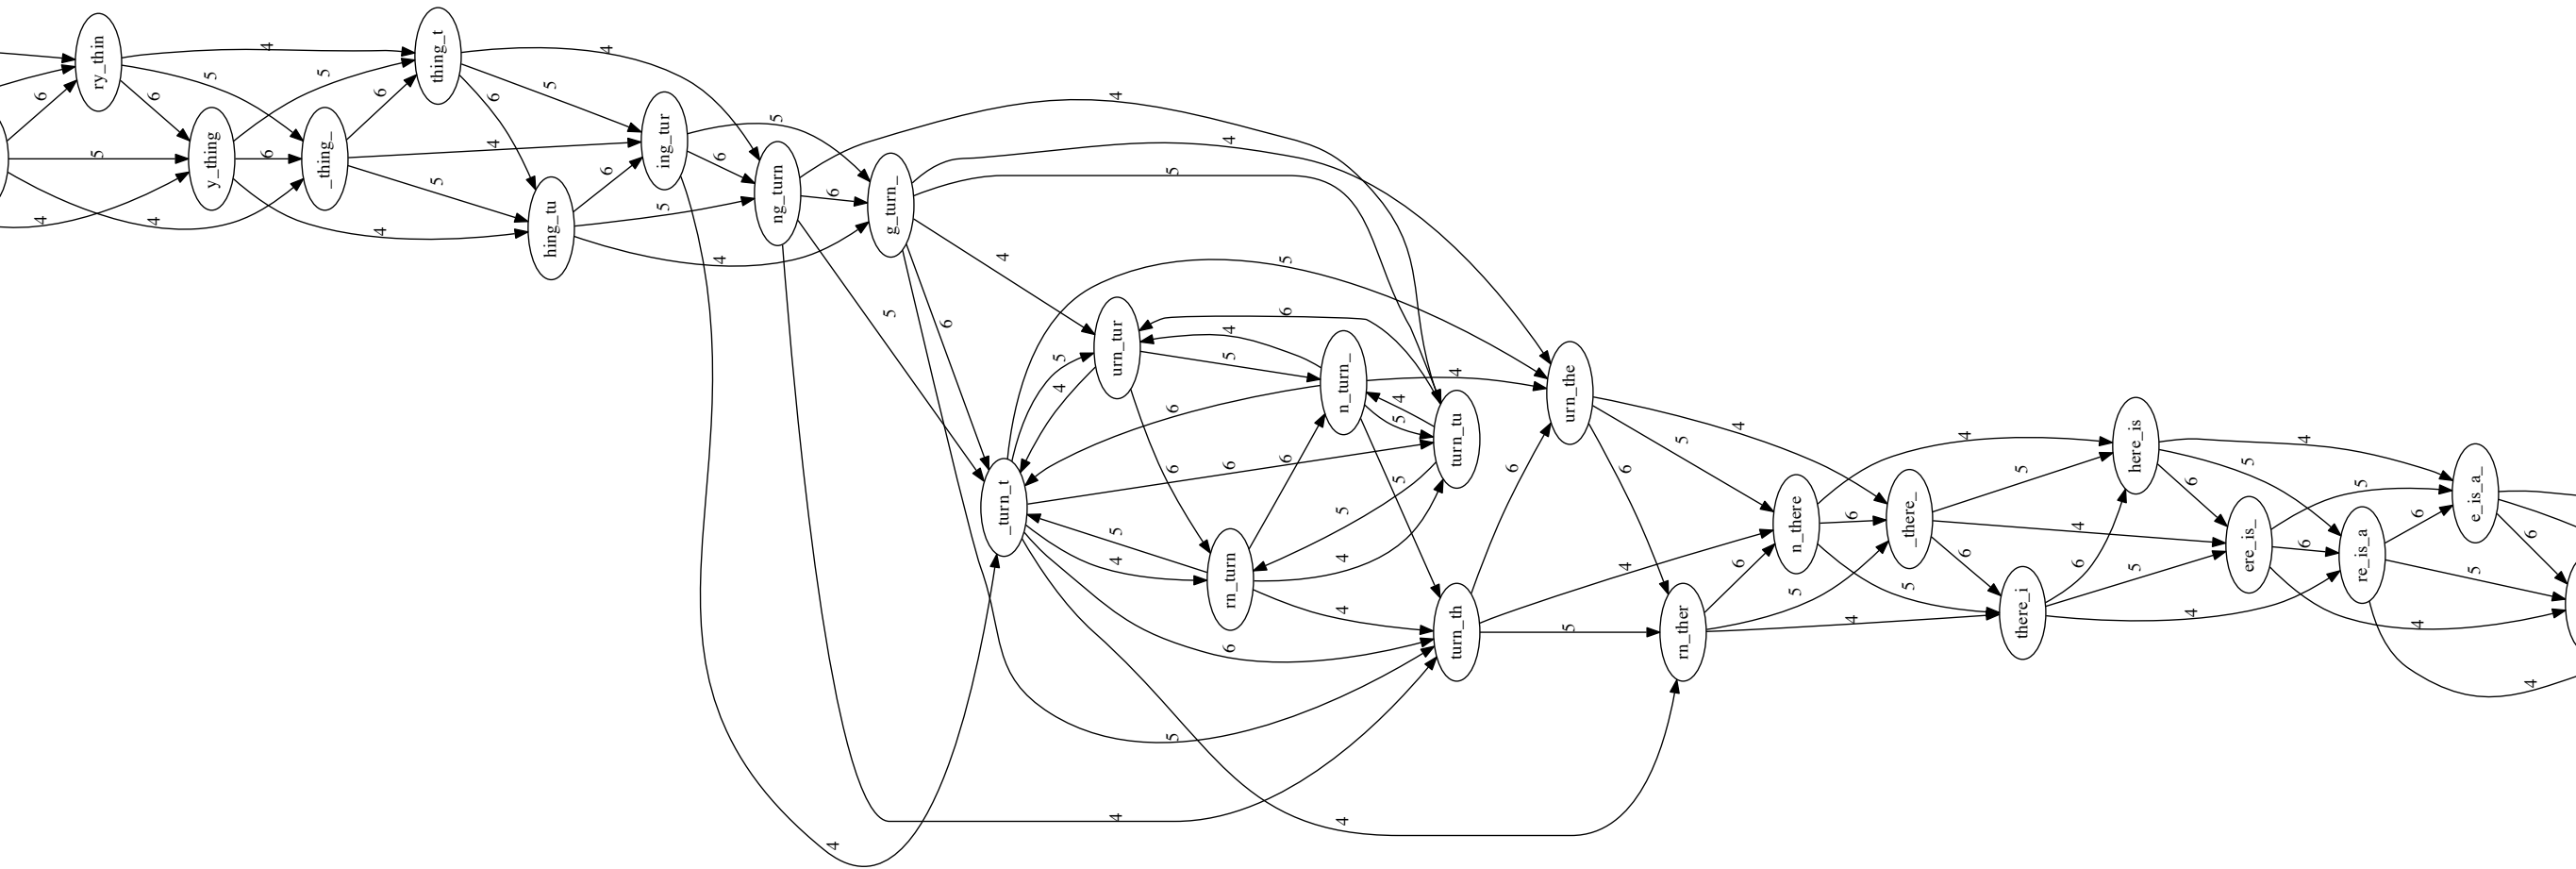


Layout

Remove transitively inferrable edges, starting with edges that skip one node:




Before:

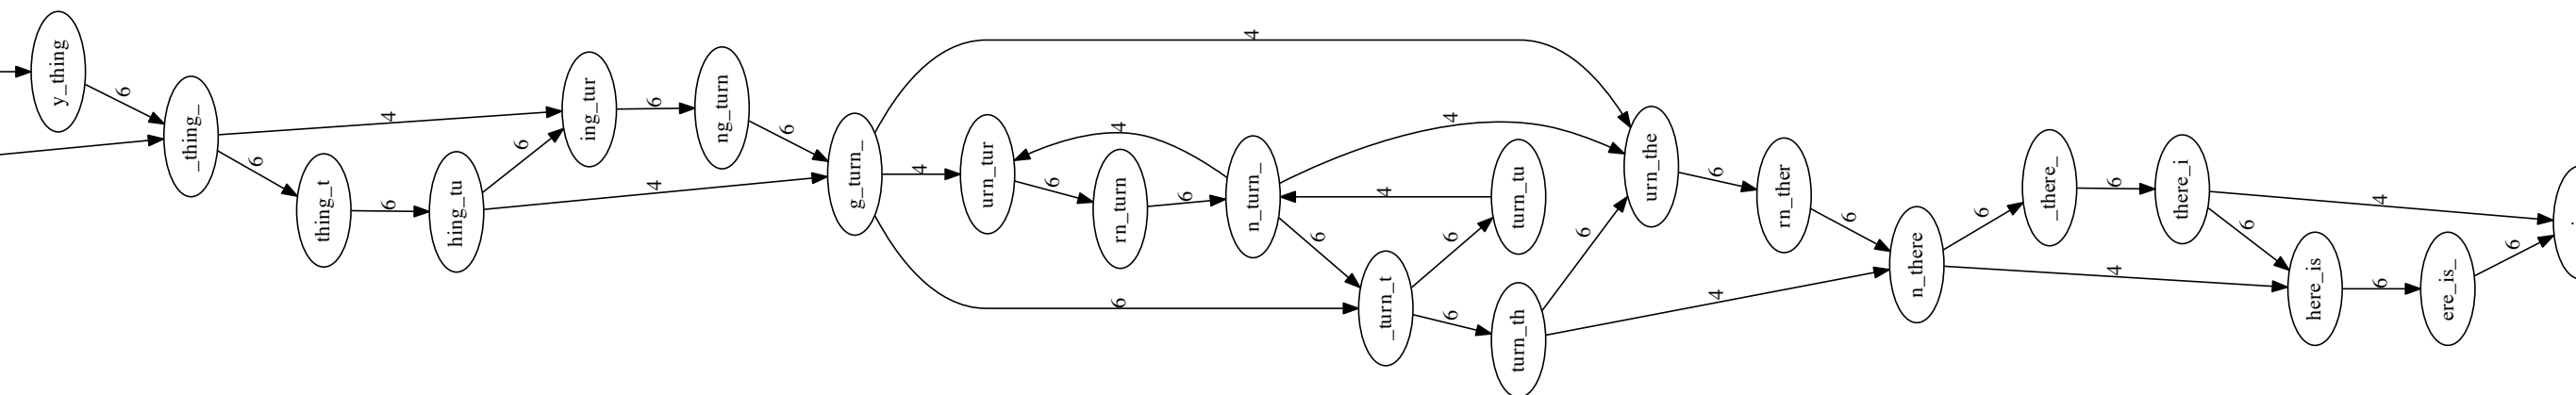


Layout

Remove transitively inferable edges, starting with edges that skip one node:




After:

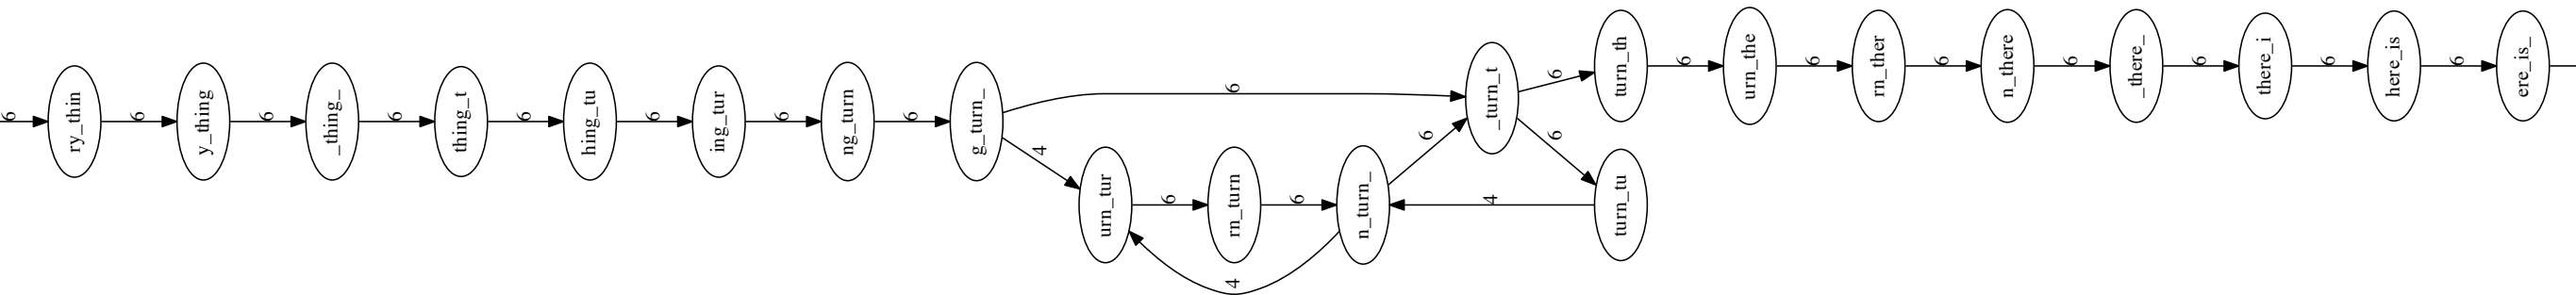


Layout

Now remove edges that skip one or two nodes:




After:



Even simpler

Layout

Emit *contigs* corresponding to the non-branching stretches



Contig 1

to_every_thing_turn_

Contig 2


turn_there_is_a_season



Unresolvable repeat

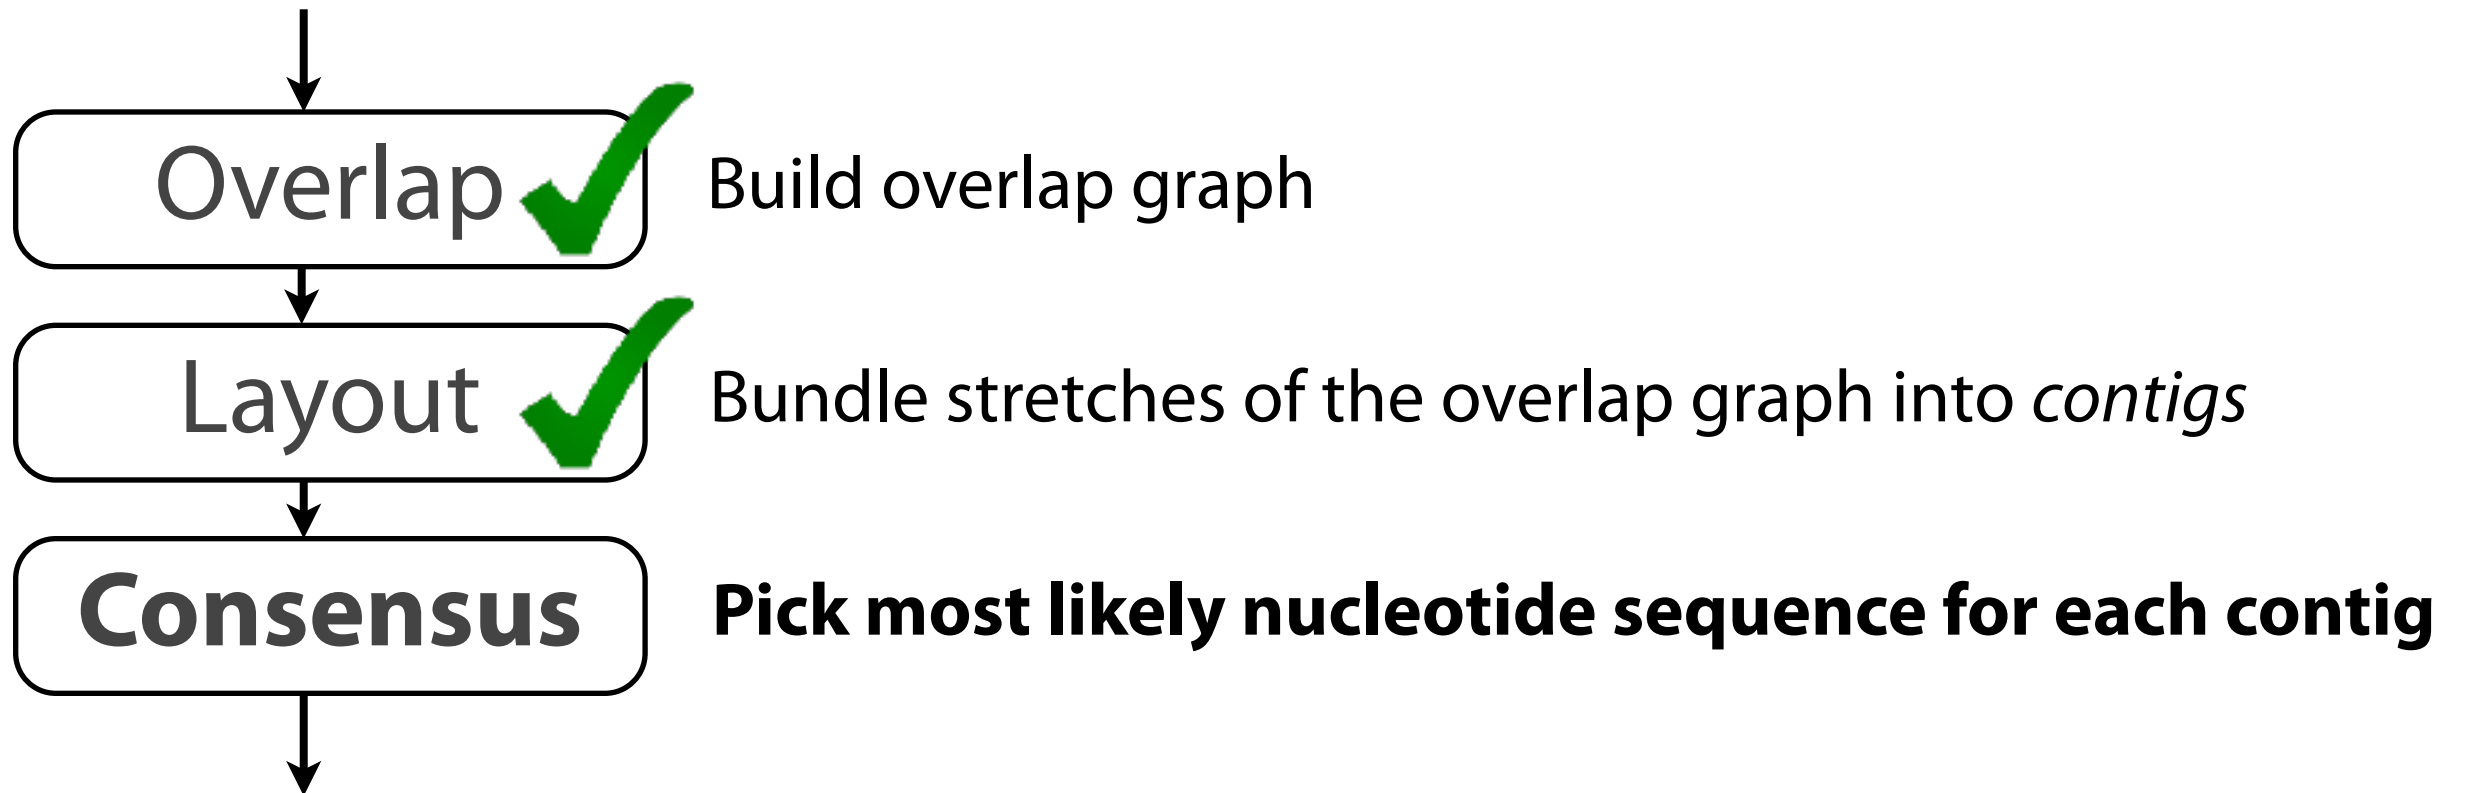
Layout

Must handle subgraphs that are spurious, e.g. because of sequencing error



Mismatch could be due to sequencing error or repeat. Since the path through **b** ends abruptly we might conclude it's an error and prune **b**.

Overlap Layout Consensus



Consensus

TAGATTACACAGATTACTGA TTGATGGCGTAA CTA
TAGATTACACAGATTACTGACTTTGATGGCGTAAACTA
TAG TTACACAGATTATTGACTTTCATGGCGTAA CTA
TAGATTACACAGATTACTGACTTTGATGGCGTAA CTA
TAGATTACACAGATTACTGACTTTGATGGCGTAA CTA

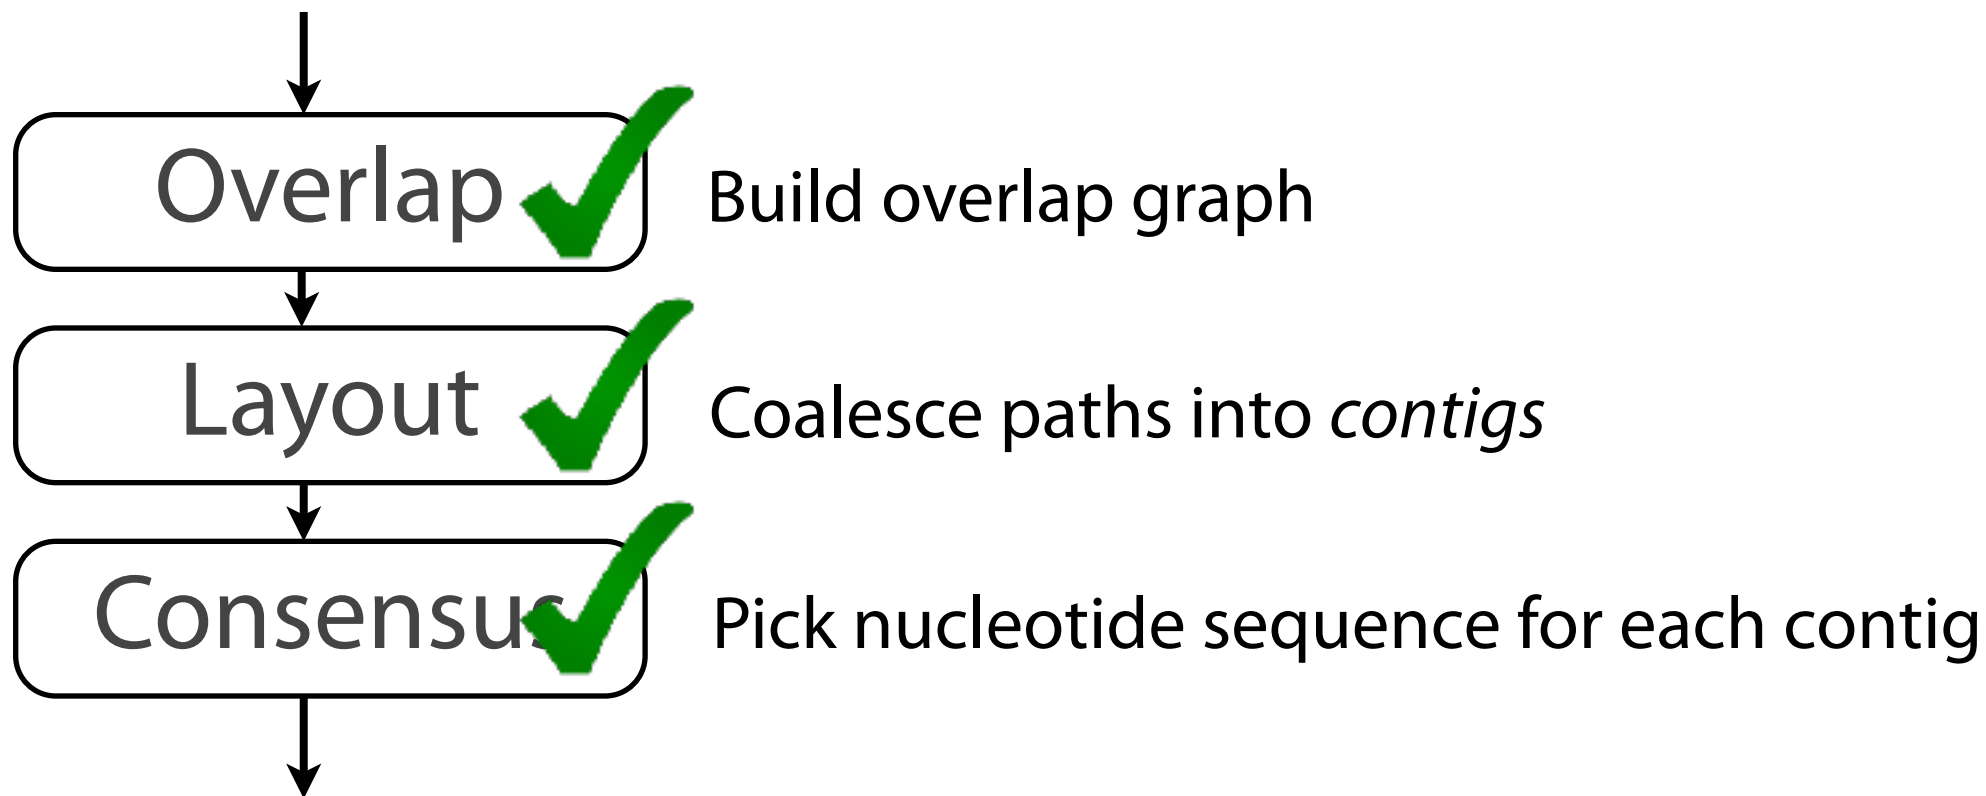
Take reads that make
up a contig and line
them up

↓ ↓ ↓ ↓ ↓
TAGATTACACAGATTACTGACTTTGATGGCGTAA CTA

Take *consensus*, i.e.
majority vote

Complications: (a) sequencing error, (b) ploidy

Overlap Layout Consensus



OLC drawbacks

Building overlap graph is slow. We saw $O(N + a)$ and $O(N^2)$ approaches.

Overlap graph is big; one node per read, # edges can grow superlinearly with # reads

Sequencing datasets are ~ 100s of millions or billions of reads